

# Intermediate *R* / *Bioconductor* for Sequence Analysis

Marc Carlson, Valerie Obenchain, Hervé Pagès, Paul Shannon, Dan Tenenbaum, Martin Morgan<sup>1</sup>

14-15 February 2013

<sup>1</sup>[mtmorgan@fhcrc.org](mailto:mtmorgan@fhcrc.org)

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
<b>I</b>	<b><i>R</i> / <i>Bioconductor</i></b>	<b>5</b>
<b>2</b>	<b><i>R</i></b>	<b>6</b>
2.1	Statistical analysis and comprehension . . . . .	6
2.2	Basics of <i>R</i> . . . . .	6
2.2.1	Essential data types . . . . .	7
2.2.2	S3 (and S4) classes . . . . .	10
2.2.3	Functions . . . . .	10
2.3	In and out of trouble . . . . .	15
2.3.1	Warnings, errors, and debugging . . . . .	15
2.3.2	Efficient <i>R</i> code . . . . .	17
2.4	Packages . . . . .	21
2.5	Help! . . . . .	23
<b>3</b>	<b><i>Bioconductor</i></b>	<b>25</b>
3.1	High-throughput sequence analysis . . . . .	26
3.2	Statistical programming . . . . .	26
3.3	<i>Bioconductor</i> packages for high-throughput sequence analysis . . . . .	28
3.4	S4 Classes and methods . . . . .	28
3.5	Help! . . . . .	31
<b>4</b>	<b>Sequencing</b>	<b>32</b>
4.1	Technologies . . . . .	32
4.2	Data . . . . .	33
4.2.1	A running example: the <i>pasilla</i> data set . . . . .	33
4.2.2	Work flows . . . . .	33
<b>5</b>	<b>Strings and Reads</b>	<b>36</b>
5.1	DNA (and other) Strings with the <i>Biostrings</i> package . . . . .	36
5.2	Reads and the <i>ShortRead</i> package . . . . .	39
<b>6</b>	<b>Ranges and Alignments</b>	<b>44</b>
6.1	Ranges and the <i>GenomicRanges</i> package . . . . .	44
6.2	Alignments and the <i>Rsamtools</i> Package . . . . .	50

<b>II</b>	<b>Differential Representation</b>	<b>55</b>
<b>7</b>	<b>RNA-seq Work Flows</b>	<b>56</b>
7.1	Varieties of RNA-seq . . . . .	56
7.2	Work flows and upstream analysis . . . . .	56
7.2.1	Experimental design . . . . .	57
7.2.2	Wet-lab protocols, sequencing, and alignment . . . . .	57
7.3	Statistical analysis . . . . .	58
7.3.1	Summarizing . . . . .	58
7.3.2	Normalization . . . . .	59
7.3.3	Error model . . . . .	59
7.3.4	Multiple comparison . . . . .	59
7.3.5	<i>Bioconductor</i> software . . . . .	60
<b>8</b>	<b>DESeq Work Flow Exercises</b>	<b>61</b>
8.1	Data input and preparation . . . . .	61
8.2	Inference . . . . .	61
8.3	Independent filtering . . . . .	62
8.4	Data quality assessment . . . . .	62
8.4.1	Preliminary transformation . . . . .	62
8.4.2	Quality assessment . . . . .	62
8.5	Frequently asked questions . . . . .	62
<b>III</b>	<b>Variant Calls</b>	<b>63</b>
<b>9</b>	<b>Variant Work Flows</b>	<b>64</b>
9.1	Variants . . . . .	64
9.1.1	Varieties of variant-related work flows . . . . .	64
9.1.2	Work flows . . . . .	64
9.1.3	<i>Bioconductor</i> software . . . . .	64
9.2	<i>VariantTools</i> . . . . .	65
9.2.1	Example data: lung cancer cell lines . . . . .	65
9.2.2	Calling single-sample variants . . . . .	66
9.2.3	Additional work flows . . . . .	66
<b>10</b>	<b>Working with Called Variants</b>	<b>67</b>
10.1	Variant call format (VCF) files with <i>VariantAnnotation</i> . . . . .	67
10.1.1	Data input . . . . .	67
10.2	SNP Annotation . . . . .	69
10.3	Large-scale filtering . . . . .	74
<b>IV</b>	<b>Annotation and Visualization</b>	<b>77</b>
<b>11</b>	<b>Gene-centric Annotation</b>	<b>78</b>
11.1	Gene-centric annotations with <i>AnnotationDbi</i> . . . . .	78
11.2	<i>biomaRt</i> and other web-based resources . . . . .	81
11.2.1	Using <i>biomaRt</i> . . . . .	81

<b>12 Genomic Annotation</b>	<b>83</b>
12.1 Whole genome sequences . . . . .	83
12.2 Gene models . . . . .	83
12.2.1 <i>TxDb.*</i> packages for model organisms . . . . .	83
12.3 UCSC tracks . . . . .	85
12.3.1 Easily creating <i>TranscriptDb</i> objects from GTF files . . . . .	85
<b>13 Visualizing Sequence Data</b>	<b>90</b>
13.1 <i>Gviz</i> . . . . .	90
13.2 <i>ggbio</i> . . . . .	91
13.3 <i>shiny</i> for easy interactive reports . . . . .	91
<b>References</b>	<b>93</b>
 <b>V Appendix</b>	 <b>95</b>
A <i>DESeq</i> vignette	96
B <i>VariantTools</i> vignette	97

# Chapter 1

## Introduction

Intermediate *R* / *Bioconductor* for High-Throughput Sequence Analysis introduces users with some *R* experience to common *Bioconductor* work flows for sequence analysis. The course involves a combination of presentations and hands-on exercises. Our starting point is BAM files created by aligning short reads to a reference genome. Topics include: exploratory analysis ([GenomicRanges](#), [Rsamtools](#)); assessing differential expression of known genes ([DESeq](#)); detection, calling, and manipulation of variants ([VariantTools](#), [VariantAnnotation](#)). We learn how to integrate results with curated gene and genomic annotations ([GenomicFeatures](#)), and to visualize results ([GViz](#), [ggbio](#)).

The course will use the ‘devel’ version of *Bioconductor*. Course participants will have access to a configured Amazon machine instance, with easy access through a web browser and *Rstudio*; the only requirement is that users have *wireless internet capabilities* and have installed a modern web browser installed<sup>1</sup>. Participants wanting to use a version of *Bioconductor* on their own laptop should come with the ‘devel’ (to be *R* 3.0.0) version of *R* installed, and should have followed the package installation instructions available from the course web page<sup>2</sup> shortly before (e.g., Wednesday morning) the start of the course. Software installation will require high-speed internet access; relevant software can be installed for take-home use during the course.

There are many books to help with using *R*, but not yet a book-length treatment of *R* / *Bioconductor* tools for sequence analysis. Starting points for bioinformatic analysis in *R*, still relevant for statistical and informatic concepts though not directly addressing sequence analysis, are Hahne’s *Bioconductor Case Studies* [8] and Gentleman’s *R Programming for Bioinformatics* [7]. For *R* novices, one place to start is Pardis’ *R for Beginners*<sup>3</sup>. General *R* programming recommendations include Dalgaard’s *Introductory Statistics with R* [6], Matloff’s *The Art of R Programming* [14], and Meys and de Vies’ *R For Dummies* [15]; an interesting internet resource for intermediate *R* programming is Burns’ *The R Inferno*<sup>4</sup>.

---

<sup>1</sup>See requirements at [http://www.rstudio.com/ide/docs/advanced/optimizing\\_browser](http://www.rstudio.com/ide/docs/advanced/optimizing_browser)

<sup>2</sup><http://bioconductor.org/help/course-materials/2013/SeattleFeb2013/>

<sup>3</sup>[http://cran.r-project.org/doc/contrib/Paradis-rdebuts\\_en.pdf](http://cran.r-project.org/doc/contrib/Paradis-rdebuts_en.pdf)

<sup>4</sup><http://www.burns-stat.com/documents/books/the-r-inferno/>

Table 1.1: Tentative schedule.

Day 1	
Morning	Part I: Up to speed with <i>R</i> scripts and <i>Bioconductor</i> packages.
Afternoon	Part II: Differential expression.
Day 2	
Morning	Part III: Variants.
Afternoon	Part IV: Annotation and Visualization.

## Part I

# *R / Bioconductor*

# Chapter 2

## *R*

### 2.1 Statistical analysis and comprehension

*R* is an open-source statistical programming language. It is used to manipulate data, to perform statistical analysis, and to present graphical and other results. *R* consists of a core language, additional ‘packages’ distributed with the *R* language, and a very large number of packages contributed by the broader community. Packages add specific functionality to an *R* installation. *R* has become the primary language of academic statistical analysis, and is widely used in diverse areas of research, government, and industry.

*R* has several unique features. It has a surprisingly ‘old school’ interface: users type commands into a console; scripts in plain text represent work flows; tools other than *R* are used for editing and other tasks. *R* is a flexible programming language, so while one person might use functions provided by *R* to accomplish advanced analytic tasks, another might implement their own functions for novel data types. As a programming language, *R* adopts syntax and grammar that differ from many other languages: objects in *R* are ‘vectors’, and functions are ‘vectorized’ to operate on all elements of the object; *R* objects have ‘copy on change’ and ‘pass by value’ semantics, reducing unexpected consequences for users at the expense of less efficient memory use; common paradigms in other languages, such as the ‘for’ loop, are encountered much less commonly in *R*. Many authors contribute to *R*, so there can be a frustrating inconsistency of documentation and interface. *R* grew up in the academic community, so authors have not shied away from trying new approaches. Common statistical analysis functions are very well-developed.

### 2.2 Basics of *R*

Opening an *R* session results in a prompt. The user types instructions at the prompt. Here is an example:

```
> ## assign values 5, 4, 3, 2, 1 to variable 'x'
> x <- c(5, 4, 3, 2, 1)
> x

[1] 5 4 3 2 1
```

The first line starts with a `#` to represent a comment; the line is ignored by *R*. The next line creates a variable `x`. The variable is assigned (using `<-`, we could have used `=` almost interchangeably) a value. The value assigned is the result of a call to the `c` function. That it is a function call is indicated by the symbol named followed by parentheses, `c()`. The `c` function takes zero or more arguments, and returns a vector. The vector is the value assigned to `x`. *R* responds to this line with a new prompt, ready for the next input. The next line asks *R* to display the value of the variable `x`. *R* responds by printing `[1]` to indicate that the subsequent number is the first element of the vector. It then prints the value of `x`.

Table 2.1: Essential aspects of the *R* language.

Category	Function	Description
Vectors	<code>integer</code> , <code>numeric</code>	Vectors holding a single type of data length 0 or more
	<code>complex</code>	<code>character</code>
	<code>raw</code>	<code>factor</code>
List-like	<code>list</code>	Arbitrary collections of elements
	<code>data.frame</code>	List of equal-length vectors
	<code>environment</code>	<i>Pass-by-reference</i> data storage
Array-like	<code>matrix</code>	Two-dimensional, homogeneous types
	<code>data.frame</code>	Homogeneous columns; row- and column indexing
	<code>array</code>	0 or more dimensions
Statistical	<code>NA</code> , <code>factor</code>	Essential statistical concepts, integral to the language.
Classes	'S3'	List-like structured data; simple inheritance & dispatch
	'S4'	Formal classes, multiple inheritance & dispatch
Functions	'function'	A simple function with arguments, body, and return value
	'generic'	A (S3 or S4) function with associated <i>methods</i>
	'method'	A function implementing a generic for an S3 or S4 class

*R* has many features to aid common operations. Entering sequences is a very common operation, and expressions of the form `2:4` create a sequence from 2 to 4. Sub-setting one vector by another is enabled with `[]`. Here we create an integer sequence from 2 to 4, and use the sequence as an index to select the second, third, and fourth elements of `x`

```
> x[2:4]
```

```
[1] 4 3 2
```

Index values can be repeated, and if outside the domain of `x` return the special value `NA`. Negative index values remove elements from the vector. Logical and character vectors (described below) can also be used for sub-setting.

*R* functions operate on variables. Functions are usually vectorized, acting on all elements of their argument and obviating the need for explicit iteration. Functions can generate warnings when performing suspect operations, or errors if evaluation cannot proceed; try `log(-1)`.

```
> log(x)
```

```
[1] 1.6094379 1.3862944 1.0986123 0.6931472 0.0000000
```

### 2.2.1 Essential data types

*R* has a number of built-in data types, summarized in Table 2.1. These represent `integer`, `numeric` (floating point), `complex`, `character`, `logical` (Boolean), and `raw` (byte) data. It is possible to convert between data types, and to discover the type or mode of a variable.

```
> c(1.1, 1.2, 1.3)          # numeric
```

```
[1] 1.1 1.2 1.3
```

```
> c(FALSE, TRUE, FALSE)    # logical
```

```
[1] FALSE TRUE FALSE
```

```
> c("foo", "bar", "baz")    # character, single or double quote ok
```



```

[1] "foo" "bar" "baz"

> as.character(x)           # convert 'x' to character

[1] "5" "4" "3" "2" "1"

> typeof(x)                 # the number 5 is numeric, not integer

[1] "double"

> typeof(2L)                # append 'L' to force integer

[1] "integer"

> typeof(2:4)               # ':' produces a sequence of integers

[1] "integer"

```

*R* includes data types particularly useful for statistical analysis, including **factor** to represent categories and **NA** (used in any vector) to represent missing values.

```

> sex <- factor(c("Male", "Female", NA), levels=c("Female", "Male"))
> sex

[1] Male    Female <NA>
Levels: Female Male

```

**Lists, data frames, and matrices** All of the vectors mentioned so far are homogeneous, consisting of a single type of element. A **list** can contain a collection of different types of elements and, like all vectors, these elements can be named to create a key-value association.

```

> lst <- list(a=1:3, b=c("foo", "bar"), c=sex)
> lst

$a
[1] 1 2 3

$b
[1] "foo" "bar"

$c
[1] Male    Female <NA>
Levels: Female Male

```

Lists can be subset like other vectors to get another list, or subset with `[[` to retrieve the actual list element; as with other vectors, sub-setting can use names

```

> lst[c(3, 1)]             # another list

$c
[1] Male    Female <NA>
Levels: Female Male

$a
[1] 1 2 3

```

```
> lst[["a"]]           # the element itself, selected by name
[1] 1 2 3
```

A `data.frame` is a list of equal-length vectors, representing a rectangular data structure not unlike a spread sheet. Each column of the data frame is a vector, so data types must be homogeneous within a column. A `data.frame` can be subset by row or column, and columns can be accessed with `$` or `[[`.

```
> df <- data.frame(age=c(27L, 32L, 19L),
+                  sex=factor(c("Male", "Female", "Male")))
> df
  age  sex
1  27 Male
2  32 Female
3  19  Male

> df[c(1, 3),]
  age sex
1  27 Male
3  19 Male

> df[df$age > 20,]
  age  sex
1  27 Male
2  32 Female
```

A `matrix` is also a rectangular data structure, but subject to the constraint that all elements are the same type. A matrix is created by taking a vector, and specifying the number of rows or columns the vector is to represent. On sub-setting, *R* coerces a single column `data.frame` or single row or column `matrix` to a vector if possible; use `drop=FALSE` to stop this behavior.

```
> m <- matrix(1:12, nrow=3)
> m
     [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12

> m[c(1, 3), c(2, 4)]
     [,1] [,2]
[1,]    4   10
[2,]    6   12

> m[, 3]
[1] 7 8 9

> m[, 3, drop=FALSE]
     [,1]
[1,]    7
[2,]    8
[3,]    9
```

An `array` is a data structure for representing homogeneous, rectangular data in higher dimensions.

### 2.2.2 S3 (and S4) classes

More complicated data structures are represented using the ‘S3’ or ‘S4’ object system. Objects are often created by functions (for example, `lm`, below), with parts of the object extracted or assigned using *accessor* functions. The following generates 1000 random normal deviates as `x`, and uses these to create another 1000 deviates `y` that are linearly related to `x` but with some error. We fit a linear regression using a ‘formula’ to describe the relationship between variables, summarize the results in a familiar ANOVA table, and access `fit` (an S3 object) for the residuals of the regression, using these as input first to the `var` (variance) and then `sqrt` (square-root) functions. Objects can be interrogated for their class.

```
> x <- rnorm(1000, sd=1)
> y <- x + rnorm(1000, sd=.5)
> fit <- lm(y ~ x)           # formula describes linear regression
> fit                        # an 'S3' object

Call:
lm(formula = y ~ x)

Coefficients:
(Intercept)          x
      0.02226       1.00869

> anova(fit)

Analysis of Variance Table

Response: y
      Df Sum Sq Mean Sq F value    Pr(>F)    
x         1 1146.99  1146.99   4651.4 < 2.2e-16 ***
Residuals 998   246.09    0.25             
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

> sqrt(var(resid(fit))) # residuals accessor and subsequent transforms

[1] 0.4963278

> class(fit)

[1] "lm"
```

Many *Bioconductor* packages implement S4 objects to represent data. S3 and S4 systems are quite different from a programmer’s perspective, but fairly similar from a user’s perspective: both systems encapsulate complicated data structures, and allow for methods specialized to different data types; accessors are used to extract information from the objects.

### 2.2.3 Functions

*R* has a very large number of functions; Table 2.2 provides a brief list of those that might be commonly used and particularly useful. See the help pages (e.g., `?lm`) and examples (`example(match)`) for more details on each of these functions.

*R* functions accept arguments, and return values. Arguments can be required or optional. Some functions may take variable numbers of arguments, e.g., the columns in a `data.frame`

Table 2.2: A selection of *R* function.

---

<code>dir</code> , <code>read.table</code> ( <b>and friends</b> ), <code>scan</code>	List files in a directory, read spreadsheet-like data into <i>R</i> , efficiently read homogeneous data (e.g., a file of numeric values) to be represented as a matrix.
<code>c</code> , <code>factor</code> , <code>data.frame</code> , <code>matrix</code>	Create a vector, factor, data frame or matrix.
<code>summary</code> , <code>table</code> , <code>xtabs</code>	Summarize, create a table of the number of times elements occur in a vector, cross-tabulate two or more variables.
<code>t.test</code> , <code>aov</code> , <code>lm</code> , <code>anova</code> , <code>chisq.test</code>	Basic comparison of two ( <code>t.test</code> ) groups, or several groups via analysis of variance / linear models ( <code>aov</code> output is probably more familiar to biologists), or compare simpler with more complicated models ( <code>anova</code> ); $\chi^2$ tests.
<code>dist</code> , <code>hclust</code>	Cluster data.
<code>plot</code>	Plot data.
<code>ls</code> , <code>str</code> , <code>library</code> , <code>search</code>	List objects in the current (or specified) workspace, or peak at the structure of an object; add a library to or describe the search path of attached packages.
<code>lapply</code> , <code>sapply</code> , <code>mapply</code> , <code>aggregate</code>	Apply a function to each element of a list ( <code>lapply</code> , <code>sapply</code> ), to elements of several lists ( <code>mapply</code> ), or to elements of a list partitioned by one or more factors ( <code>aggregate</code> ).
<code>with</code>	Conveniently access columns of a data frame or other element without having to repeat the name of the data frame.
<code>match</code> , <code>%in%</code>	Report the index or existence of elements from one vector that match another.
<code>split</code> , <code>cut</code> , <code>unlist</code>	Split one vector by an equal length factor, cut a single vector into intervals encoded as levels of a factor, <code>unlist</code> (concatenate) list elements.
<code>strsplit</code> , <code>grep</code> , <code>sub</code>	Operate on character vectors, splitting it into distinct fields, searching for the occurrence of a patterns using regular expressions (see <code>?regex</code> , or substituting a string for a regular expression.
<code>install.packages</code>	Install a package from an on-line repository into your <i>R</i> .
<code>traceback</code> , <code>debug</code> , <code>browser</code>	Report the sequence of functions under evaluation at the time of the error; enter a debugger when a particular function or statement is invoked.

---

```

> y <- 5:1
> log(y)

[1] 1.6094379 1.3862944 1.0986123 0.6931472 0.0000000

> args(log)          # arguments 'x' and 'base'; see ?log

function (x, base = exp(1))
NULL

> log(y, base=2)     # 'base' is optional, with default value

[1] 2.321928 2.000000 1.584963 1.000000 0.000000

> try(log())         # 'x' required; 'try' continues even on error
> args(data.frame)  # ... represents variable number of arguments

function (... , row.names = NULL, check.rows = FALSE, check.names = TRUE,
  stringsAsFactors = default.stringsAsFactors())
NULL

```

Arguments can be matched by name or position. If an argument appears after ..., it must be named.

```

> log(base=2, y)     # match argument 'base' by name, 'x' by position

[1] 2.321928 2.000000 1.584963 1.000000 0.000000

```

A function such as `anova` is a *generic* that provides an overall signature but dispatches the actual work to the *method* corresponding to the class(es) of the arguments used to invoke the generic. A generic may have fewer arguments than a method, as with the S3 function `anova` and its method `anova.glm`.

```

> args(anova)

function (object, ...)
NULL

> args(anova.glm)

function (object, ..., dispersion = NULL, test = NULL)
NULL

```

The ... argument in the `anova` generic means that additional arguments are possible; the `anova` generic hands these arguments to the method it dispatches to.

### Exercise 1

This exercise uses data describing 128 microarray samples as a basis for exploring R functions. Covariates such as age, sex, type, stage of the disease, etc., are in a data file `pData.csv`.

The following command creates a variable `pdataFiles` that is the location of a comma-separated value ('csv') file to be used in the exercise. A csv file can be created using, e.g., 'Save as...' in spreadsheet software.

```

> pdataFile <- system.file(package="SequenceAnalysisData", "extdata",
+                           "pData.csv")

```

Input the csv file using `read.table`, assigning the input to a variable `pdata`. Use `dim` to find out the dimensions (number of rows, number of columns) in the object. Are there 128 rows? Use `names` or `colnames` to list the names of the columns of `pdata`. Use `summary` to summarize each column of the data. What are the data types of each column in the data frame?

A data frame is a list of equal length vectors. Select the 'sex' column of the data frame using `[[` or `$`. Pause to explain to your neighbor why this sub-setting works. Since a data frame is a list, use `sapply` to ask about the class of each column in the data frame. Explain to your neighbor what this produces, and why.

Use `table` to summarize the number of males and females in the sample. Consult the help page `?table` to figure out additional arguments required to include NA values in the tabulation.

The `mol.biol` column summarizes molecular biological attributes of each sample. Use `table` to summarize the different molecular biology levels in the sample. Use `%in%` to create a logical vector of the samples that are either BCR/ABL or NEG. Subset the original phenotypic data to contain those samples that are BCR/ABL or NEG.

After sub-setting, what are the levels of the `mol.biol` column? Set the levels to be BCR/ABL and NEG, i.e., the levels in the subset.

One would like covariates to be similar across groups of interest. Use `t.test` to assess whether BCR/ABL and NEG have individuals with similar age. To do this, use a formula that describes the response age in terms of the predictor `mol.biol`. If age is not independent of molecular biology, what complications might this introduce into subsequent analysis? Use

**Solution:** Here we input the data and explore basic properties.

```
> pdata <- read.table(pdataFile)
> dim(pdata)

[1] 128 21

> names(pdata)

[1] "cod"           "diagnosis"      "sex"            "age"
[5] "BT"           "remission"      "CR"             "date.cr"
[9] "t.4.11."       "t.9.22."        "cyto.normal"    "citog"
[13] "mol.biol"      "fusion.protein" "mdr"            "kinet"
[17] "ccr"           "relapse"        "transplant"     "f.u"
[21] "date.last.seen"

> summary(pdata)

      cod      diagnosis      sex      age      BT
10005 : 1 11/15/1997: 2 F :42 Min. : 5.00 B2 :36
1003 : 1 1/15/1997 : 2 M :83 1st Qu.:19.00 B3 :23
remission      CR      date.cr      t.4.11.
CR :99 CR :96 11/11/1997: 3 Mode :logical
REF :15 DEATH IN CR : 3 10/18/1999: 2 FALSE:86
t.9.22.      cyto.normal      citog      mol.biol
Mode :logical Mode :logical normal :24 ALL1/AF4:10
FALSE:67 FALSE:69 simple alt. :15 BCR/ABL :37
fusion.protein mdr      kinet      ccr      relapse
p190 :17 NEG :101 dyplloid:94 Mode :logical Mode :logical
p190/p210: 8 POS : 24 hyperd.:27 FALSE:74 FALSE:35
transplant      f.u      date.last.seen
Mode :logical REL :61 12/15/1997: 2
FALSE:91 CCR :23 12/31/2002: 2
[ reached getOption("max.print") -- omitted 5 rows ]
```

A data frame can be subset as if it were a matrix, or a list of column vectors.

```
> head(pdata[, "sex"], 3)
```

```
[1] M M F
Levels: F M
```

```
> head(pdata$sex, 3)
```

```
[1] M M F
Levels: F M
```

```
> head(pdata[["sex"]], 3)
```

```
[1] M M F
Levels: F M
```

```
> sapply(pdata, class)
```

```
      cod      diagnosis      sex      age      BT
"factor"  "factor"    "factor"  "integer" "factor"
remission      CR      date.cr    t.4.11.    t.9.22.
"factor"  "factor"    "factor"  "logical" "logical"
cyto.normal      citog  mol.biol fusion.protein      mdr
"logical"  "factor"    "factor"  "factor"  "factor"
kinet      ccr      relapse  transplant      f.u
"factor"  "logical"  "logical"  "logical"  "factor"
date.last.seen
"factor"
```

The number of males and females, including NA, is

```
> table(pdata$sex, useNA="ifany")
```

```
  F    M <NA>
42   83    3
```

An alternative version of this uses the `with` function: `with(pdata, table(sex, useNA="ifany"))`.

The `mol.biol` column contains the following samples:

```
> with(pdata, table(mol.biol, useNA="ifany"))
```

```
mol.biol
ALL1/AF4 BCR/ABL E2A/PBX1      NEG  NUP-98  p15/p16
      10      37      5      74      1      1
```

A logical vector indicating that the corresponding row is either `BCR/ABL` or `NEG` is constructed as

```
> ridx <- pdata$mol.biol %in% c("BCR/ABL", "NEG")
```

We can get a sense of the number of rows selected via `table` or `sum` (discuss with your neighbor what `sum` does, and why the answer is the same as the number of `TRUE` values in the result of the `table` function).

```
> table(ridx)
```

```
ridx
FALSE  TRUE
   17   111
```

```
> sum(ridx)
```

```
[1] 111
```

The original data frame can be subset to contain only BCR/ABL or NEG samples using the logical vector `ridx` that we created.

```
> pdata1 <- pdata[ridx,]
```

The levels of each factor reflect the levels in the original object, rather than the levels in the subset object, e.g.,

```
> levels(pdata1$mol.biol)
```

```
[1] "ALL1/AF4" "BCR/ABL" "E2A/PBX1" "NEG"      "NUP-98"  "p15/p16"
```

These can be re-coded by updating the new data frame to contain a factor with the desired levels.

```
> pdata1$mol.biol <- factor(pdata1$mol.biol)
```

```
> table(pdata1$mol.biol)
```

BCR/ABL	NEG
37	74

To ask whether age differs between molecular biology types, we use a formula `age ~ mol.biol` to describe the relationship ('age as a function of molecular biology') that we wish to test

```
> with(pdata1, t.test(age ~ mol.biol))
```

Welch Two Sample t-test

data: age by mol.biol

t = 4.8172, df = 68.529, p-value = 8.401e-06

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

7.13507 17.22408

sample estimates:

mean in group BCR/ABL	mean in group NEG
40.25000	28.07042

This summary can be visualize with, e.g., the `boxplot` function

```
> ## not evaluated
```

```
> boxplot(age ~ mol.biol, pdata1)
```

Molecular biology seem to be strongly associated with age; individuals in the **NEG** group are considerably younger than those in the **BCR/ABL** group. We might wish to include age as a covariate in any subsequent analysis seeking to relate molecular biology to gene expression.

## 2.3 In and out of trouble

### 2.3.1 Warnings, errors, and debugging

*R* signals unexpected results through warnings and errors. Warnings occur when the calculation produces an unusual result that nonetheless does not preclude further evaluation. For instance `log(-1)` results in a value `NaN` ('not a number') that allows computation to continue, but at the same time signals an warning



Table 2.3: Tools for debugging and error-handling.

Function	Description
<code>traceback</code>	Report the ‘call stack’ at the time of an error.
<code>options(error=)</code>	Set a handler to be executed on error, e.g., <code>error=recover</code> .
<code>debug, trace</code>	Enter the browser when a function is called
<code>browser</code>	Interactive debugging
<code>tryCatch</code>	Handle an error condition in a script.
<code>system.time</code>	Time required to evaluate an expression
<code>Rprof</code>	Time spent in each function; also <code>summaryRprof</code>
<code>tracemem</code>	Indicate when memory copies occur ( <i>R</i> must be configured to support this)

```
> log(-1)
[1] NaN
Warning message:
In log(-1) : NaNs produced
```

Errors result when the inputs or outputs of a function are such that no further action can be taken, e.g., trying to take the square root of a character vector

```
> sqrt("two")
Error in sqrt("two") : Non-numeric argument to mathematical function
```

Warnings and errors occurring at the command prompt are usually easy to diagnose. They can be more enigmatic when occurring in a function, and exacerbated by sometimes cryptic (when read out of context) error messages. Some key tools for figuring out (‘debugging’) errors are summarized in Table 2.3.

An initial step in coming to terms with errors is to simplify the problem as much as possible, aiming for a ‘reproducible’ error. The reproducible error might involve a very small (even trivial) data set that immediately provokes the error. Often the process of creating a reproducible example helps to clarify what the error is, and what possible solutions might be.

Invoking `traceback()` immediately after an error occurs provides a ‘stack’ of the function calls that were in effect when the error occurred. This can help understand the context in which the error occurred. Knowing the context, one might use `debug` (or its more elaborate cousin, `trace`) to enter into a browser (see `?browser`) that allows one to step through the function in which the error occurred.

It can sometimes be useful to use global options (see `?options`) to influence what happens when an error occurs. Two common global options are `error` and `warn`. Setting `error=recover` combines the functionality of `traceback` and `debug`, allowing the user to enter the browser at any level of the call stack in effect at the time the error occurred. Default error behavior can be restored with `options(error=NULL)`. Setting `warn=2` causes warnings to be promoted to errors. For instance, initial investigation of an error might show that the error occurs when one of the arguments to a function has value `NaN`. The error might be accompanied by a warning message that the `NaN` has been introduced, but because warnings are by default not reported immediately it is not clear where the `NaN` comes from. `warn=2` means that the warning is treated as an error, and hence can be debugged using `traceback`, `debug`, and so on.

It is possible to continue evaluation even after an error occurs. The simplest mechanism uses the `try` function, but an only slightly more complicated version providing greater flexibility is `tryCatch`. `tryCatch` allows one to write a `handler` (the `error` argument to `tryCatch`, below) to address common faults in a way that allows a script to continue executing. Suppose a function `f` fails under certain conditions

```
> f <- function(i) {
+   if (i < 0)
+     stop("i is negative")
+   rnorm(i)
+ }
```

Table 2.4: Common ways to improve efficiency of *R* code.

Easy	Moderate
1. Selective input	1. Know relevant packages
2. Vectorize	2. Understand algorithm complexity
3. Pre-allocate and fill	3. Use parallel evaluation
4. Avoid expensive conveniences	4. Exploit libraries and C++ code

```
+ }
> lapply(0:1, f)
```

```
[[1]]
numeric(0)
```

```
[[2]]
[1] -0.416157
```

but we wish to continue, e.g., replacing failed conditions with NA:

```
> lapply(-1:1, function(i) {
+   tryCatch({
+     f(i)
+   }, error=function(err) {
+     ## return 'NA' when error occurs, instead of stopping
+     NA_real_
+   })
+ })
```

```
[[1]]
[1] NA
```

```
[[2]]
numeric(0)
```

```
[[3]]
[1] 0.8123937
```

### 2.3.2 Efficient *R* code

There are often many ways to accomplish a result in *R*, but these different ways often have very different speed or memory requirements. For small data sets these performance differences are not that important, but for large data sets (e.g., high-throughput sequencing; genome-wide association studies, GWAS) or complicated calculations (e.g., bootstrapping) performance can be important. Several approaches to achieving efficient *R* programming are summarized in Table 2.4.

**Easy solutions** Several common performance bottlenecks often have easy solutions; these are outlined here.

Text files often contain more information, for example 1000's of individuals at millions of SNPs, when only a subset of the data is required, e.g., during algorithm development. Reading in all the data can be demanding in terms of both memory and time. A solution is to use arguments such as `colClasses` to specify

the columns and their data types that are required, and to use `nrow` to limit the number of rows input. For example, the following ignores the first and fourth column, reading in only the second and third (as type `integer` and `numeric`).

```
> ## not evaluated
> colClasses <-
+   c("NULL", "integer", "numeric", "NULL")
> df <- read.table("myfile", colClasses=colClasses)
```

R is vectorized, so traditional programming for loops are often not necessary. Rather than calculating 100000 random numbers one at a time, or squaring each element of a vector, or iterating over rows and columns in a matrix to calculate row sums, invoke the single function that performs each of these operations.

```
> x <- runif(100000); x2 <- x^2
> m <- matrix(x2, nrow=1000); y <- rowSums(m)
```

This often requires a change of thinking, turning the sequence of operations ‘inside-out’. For instance, calculate the log of the square of each element of a vector by calculating the square of all elements, followed by the log of all elements `x2 <- x^2; x3 <- log(x2)`, or simply `logx2 <- log(x^2)`.

It may sometimes be natural to formulate a problem as a `for` loop, or the formulation of the problem may require that a `for` loop be used. In these circumstances the appropriate strategy is to pre-allocate the `result` object, and to fill the result in during loop iteration.

```
> ## not evaluated
> result <- numeric(nrow(df))
> for (i in seq_len(nrow(df)))
+   result[[i]] <- some_calc(df[i,])
```

Some R operations are helpful in general, but misleading or inefficient in particular circumstances. An example is the behavior of `unlist` when the list is named – R creates new names that have been made unique. This can be confusing (e.g., when Entrez gene identifiers are ‘mangled’ to unintentionally look like other identifiers) and expensive (when a large number of new names need to be created). Avoid creating unnecessary names, e.g.,

```
> unlist(list(a=1:2)) # name 'a' becomes 'a1', 'a2'

a1 a2
1  2

> unlist(list(a=1:2), use.names=FALSE) # no names

[1] 1 2
```

Names can be very useful for avoiding book-keeping errors, but are inefficient for repeated look-ups; use vectorized access or numeric indexing.

**Moderate solutions** Several solutions to inefficient code require greater knowledge to implement.

Using appropriate functions can greatly influence performance; it takes experience to know when an appropriate function exists. For instance, the `lm` function could be used to assess differential expression of each gene on a microarray, but the *limma* package implements this operation in a way that takes advantage of the experimental design that is common to each probe on the microarray, and does so in a very efficient manner.

```
> ## not evaluated
> library(limma) # microarray linear models
> fit <- lmFit(eSet, design)
```

Using appropriate algorithms can have significant performance benefits, especially as data becomes larger. This solution requires moderate skills, because one has to be able to think about the complexity (e.g., expected number of operations) of an algorithm, and to identify algorithms that accomplish the same goal in fewer steps. For example, a naive way of identifying which of 100 numbers are in a set of size 10 might look at all  $100 \times 10$  combinations of numbers (i.e., polynomial time), but a faster way is to create a ‘hash’ table of one of the set of elements and probe that for each of the other elements (i.e., linear time). The latter strategy is illustrated with

```
> x <- 1:100; s <- sample(x, 10)
> inS <- x %in% s
```

Parallel evaluation on several cores of a single Linux or MacOS computer is particularly easy to achieve when the code is already vectorized. The solution on these operating systems is to use the *parallel* package (part of the base *R* distribution) with functions `mclapply` or `pvec`. These functions allow the ‘master’ process to ‘fork’ processes for parallel evaluation on each of the cores of a single machine. The forked processes initially share memory with the master process, and only make copies when the forked process makes a copy (‘copy on change’ semantics). The *parallel* package does not support fork-like behavior on windows, where users need to more explicitly create a cluster of *R* workers and arrange for each to have the same data loaded into memory; similarly, parallel evaluation across computers (e.g., in a cluster) require more elaborate efforts to coordinate workers; this is typically done using `lapply`-like functions provided by the *parallel* package but specialized for simple (‘snow’) or more robust (‘MPI’) communication protocols between workers. On Linux and MacOS, the `mclapply` function is meant to be a ‘drop-in’ replacement for `lapply`, but with iterations being evaluated on different cores. The following illustrates their use, for a toy vectorized function `f`:

```
> library(parallel)
> f <- function(i) {
+   cat("'f' called, length(i) = ", length(i), "\n")
+   sqrt(i)
+ }
> res0 <- mclapply(1:5, f, mc.cores=2)

'f' called, length(i) = 1
'f' called, length(i) = 1
'f' called, length(i) = 1
'f' called, length(i) = 1
'f' called, length(i) = 1

> res1 <- pvec(1:5, f, mc.cores=2)

'f' called, length(i) = 3
'f' called, length(i) = 2

> identical(unlist(res0), res1)

[1] TRUE
```

`pvec` takes a vectorized function and distributes computation of different chunks of the vector across cores. Both functions allow the user to specify the number of cores used, and how the data are divided into chunks. It is not uncommon for execution time to scale approximately inversely with the number of available cores. There are several areas that require attention. The total amount of memory available on a single computer is fixed, so one usually wants to *iterate* through large data in chunks. Random numbers need to be synchronized across cores to avoid generating the same sequences on each ‘independent’ computation. Moving data to and especially from cores to the manager can be expensive, so strategies that minimize explicit movement (e.g., passing file names data base queries rather than *R* objects read from files; reducing data on the worker before

transmitting results to the manager) can be important. Relevant resources include the *parallel* vignette [19] and a now-dated review [24].

R is an interpreted language, and for very challenging computational problems it may be appropriate to write critical stages of an analysis in a compiled language like C or Fortran, or to use an existing programming library (e.g., the **BOOST** library) that efficiently implements advanced algorithms. R has a well-developed interface to C or Fortran, so it is ‘easy’ to do this; the *Rcpp* package provides a very nice approach for those with a little familiarity with C++ concepts. This places a significant burden on the person implementing the solution, requiring knowledge of two or more computer languages and of the interface between them.

**Measuring performance** When trying to improve performance, one wants to ensure (a) that the new code is actually faster than the previous code, and (b) both solutions arrive at the same, correct, answer.

The `system.time` function is a straight-forward way to measure the length of time a portion of code takes to evaluate.

```
> m <- matrix(runif(200000), 20000)
> system.time(apply(m, 1, sum))
```

```
      user  system elapsed
0.264    0.016    0.281
```

When comparing performance of different functions, it is appropriate to replicate timings to average over vagaries of system use, and to shuffle the order in which timings of alternative algorithms are calculated to avoid artifacts such as initial memory allocation. Rather than creating *ad hoc* approaches to timing, it is convenient to use packages such as *rbenchmark*:

```
> library(rbenchmark)
> f0 <- function(x) apply(x, 1, sum)
> f1 <- function(x) rowSums(x)
> benchmark(f0(m), f1(m),
+           columns=c("test", "elapsed", "relative"),
+           replications=5)
```

```
      test elapsed relative
1 f0(m)    1.423  203.286
2 f1(m)    0.007    1.000
```

Speed is an important metric, but equivalent results are also needed. The functions `identical` and `all.equal` provide different levels of assessing equivalence, with `all.equal` providing ability to ignore some differences, e.g., in the names of vector elements.

```
> res1 <- apply(m, 1, sum)
> res2 <- rowSums(m)
> identical(res1, res2)
```

```
[1] TRUE
```

```
> identical(c(1, -1), c(x=1, y=-1))
```

```
[1] FALSE
```

```
> all.equal(c(1, -1), c(x=1, y=-1),
+           check.attributes=FALSE)
```

```
[1] TRUE
```

Table 2.5: Selected base and contributed packages.

Package	Description
<i>base</i>	Data input and manipulation; scripting and programming.
<i>stats</i>	Essential statistical and plotting functions.
<i>lattice</i> , <i>ggplot2</i>	Approaches to advanced graphics.
<i>methods</i>	‘S4’ classes and methods.
<i>parallel</i>	Facilities for parallel evaluation.

Two additional functions for assessing performance are `Rprof` and `tracemem`; these are mentioned only briefly here. The `Rprof` function profiles *R* code, presenting a summary of the time spent in each part of several lines of *R* code. It is useful for gaining insight into the location of performance bottlenecks when these are not readily apparent from direct inspection. Memory management, especially copying large objects, can frequently contribute to poor performance. The `tracemem` function allows one to gain insight into how *R* manages memory; insights from this kind of analysis can sometimes be useful in restructuring code into a more efficient sequence.

## 2.4 Packages

Packages provide functionality beyond that available in base *R*. There are over 4000 packages in CRAN (comprehensive *R* archive network) and more than 600 *Bioconductor* packages. Packages are contributed by diverse members of the community; they vary in quality (many are excellent) and sometimes contain idiosyncratic aspects to their implementation. Table 2.5 outlines key base packages and selected contributed packages; see a local CRAN mirror (including the [task views](#) summarizing packages in different domains) and *Bioconductor* for additional contributed packages.

The *lattice* package illustrates the value packages add to base *R*. *lattice* is distributed with *R* but not loaded by default. It provides a very expressive way to visualize data. The following example plots yield for a number of barley varieties, conditioned on site and grouped by year. Figure 2.1 is read from the lower left corner. Note the common scales, efficient use of space, and not-too-pleasing default color palette. The Morris sample appears to be mis-labeled for ‘year’, an apparent error in the original data. Find out about the built-in data set used in this example with `?barley`.

```
> library(lattice)
> plt <- dotplot(variety ~ yield | site, data = barley, groups = year,
+               xlab = "Barley Yield (bushels/acre)" , ylab=NULL,
+               key = simpleKey(levels(barley$year), space = "top",
+               columns=2),
+               aspect=0.5, layout = c(2,3))
> print(plt)
```

New packages can be added to an *R* installation using `install.packages`. A package is installed only once per *R* installation, but needs to be loaded (with `library`) in each session in which it is used. Loading a package also loads any package that it depends on. Packages loaded in the current session are displayed with `search`. The ordering of packages returned by `search` represents the order in which the global environment (where commands entered at the prompt are evaluated) and attached packages are searched for symbols; it is possible for a package earlier in the search path to mask symbols later in the search path; these can be disambiguated using `::`.

```
> length(search())
```

```
[1] 22
```




Figure 2.1: Variety yield conditional on site and grouped by year, for the `barley` data set. What's up with Morris?

```
> search()

[1] ".GlobalEnv"                "package:lattice"
[3] "package:rbenchmark"        "package:SequenceAnalysisData"
[5] "package:edgeR"              "package:limma"
[7] "package:GenomicFeatures"    "package:AnnotationDbi"
[9] "package:Biobase"            "package:GenomicRanges"
[11] "package:IRanges"            "package:BiocGenerics"
[13] "package:parallel"           "package:BiocInstaller"
[15] "package:stats"              "package:graphics"
[17] "package:grDevices"          "package:utils"
[19] "package:datasets"           "package:methods"
[21] "Autoloads"                  "package:base"

> base::log(1:3)

[1] 0.0000000 0.6931472 1.0986123
```

## Exercise 2

Use the `library` function to load the `IntermediateSequenceAnalysis2013` package. Use the `sessionInfo` function to verify that you are using R version 2.15.2 and current packages, similar to those reported here. What other packages were loaded along with `IntermediateSequenceAnalysis2013`?

## Solution:

```
> library(IntermediateSequenceAnalysis2013)
> sessionInfo()
```

## 2.5 Help!

Find help using the *R* help system. Start a web browser with

```
> help.start()
```

The ‘Search Engine and Keywords’ link is helpful in day-to-day use.

**Manual pages** Use manual pages to find detailed descriptions of the arguments and return values of functions, and the structure and methods of classes. Find help within an *R* session as

```
> ?data.frame
> ?lm
> ?anova          # a generic function
> ?anova.lm       # an S3 method, specialized for 'lm' objects
```

S3 methods can be queried interactively. For S3,

```
> methods(anova)

[1] anova.glm      anova.glmlist anova.lm      anova.loess*  anova.MAList
[6] anova.mlm      anova.nls*
```

Non-visible functions are asterisked

```
> methods(class="glm")

[1] add1.glm*      anova.glm      confint.glm*
[4] cooks.distance.glm* deviance.glm*  drop1.glm*
[7] effects.glm*   extractAIC.glm* family.glm*
[10] formula.glm*  influence.glm* logLik.glm*
[13] model.frame.glm  nobs.glm*     predict.glm
[16] print.glm      residuals.glm  rstandard.glm
[19] rstudent.glm   summary.glm    vcov.glm*
[22] weights.glm*
```

Non-visible functions are asterisked

It is often useful to view a method definition, either by typing the method name at the command line or, for ‘non-visible’ methods, using `getAnywhere`:

```
> anova.lm
> getAnywhere("anova.loess")
```

For instance, the source code of a function is printed if the function is invoked without parentheses. Here we discover that the function `head` (which returns the first 6 elements of anything) defined in the *utils* package, is an S3 generic (indicated by `UseMethod`) and has several methods. We use `head` to look at the first six lines of the `head` method specialized for *matrix* objects.

```
> utils::head

function (x, ...)
UseMethod("head")
<environment: namespace:utils>

> methods(head)
```



```
[1] head.data.frame* head.default*   head.ftable*   head.function*
[5] head.matrix      head.table*
```

Non-visible functions are asterisked

```
> head(head.matrix)

1 function (x, n = 6L, ...)
2 {
3     stopifnot(length(n) == 1L)
4     n <- if (n < 0L)
5         max(nrow(x) + n, 0L)
6     else min(n, nrow(x))
```

**Vignettes** Vignettes, especially in *Bioconductor* packages, provide an extensive narrative describing overall package functionality. Use

```
> vignette(package="IntermediateSequenceAnalysis2013")
```

to see, in your web browser, vignettes available in the *IntermediateSequenceAnalysis2013* package. Vignettes usually consist of text with embedded *R* code, a form of literate programming. The vignette can be read as a PDF document, while the *R* source code is present as a script file ending with extension `.R`. The script file can be sourced or copied into an *R* session to evaluate exactly the commands used in the vignette. For *Bioconductor* packages, vignettes are available on the package ‘landing page’, e.g., for *IRanges*<sup>1</sup>

---

<sup>1</sup><http://bioconductor.org/packages/devel/bioc/html/IRanges.html>

## Chapter 3

# Bioconductor

*Bioconductor* is a collection of *R* packages for the analysis and comprehension of high-throughput genomic data. *Bioconductor* started more than 10 years ago. It gained credibility for its statistically rigorous approach to microarray pre-processing and analysis of designed experiments, and integrative and reproducible approaches to bioinformatic tasks. There are now more than 600 *Bioconductor* packages for expression and other microarrays, sequence analysis, flow cytometry, imaging, and other domains. The *Bioconductor* [web site](http://www.bioconductor.org) provides installation, package repository, help, and other documentation.

The *Bioconductor* web site is at [bioconductor.org](http://www.bioconductor.org). Features include:

- Introductory [work flows](#).
- A manifest of [Bioconductor packages](#) arranged in [BiocViews](#).
- [Annotation](#) (data bases of relevant genomic information, e.g., Entrez gene ids in model organisms, KEGG pathways) and [experiment data](#) (containing relatively comprehensive data sets and their analysis) packages.
- [Mailing lists](#), including searchable archives, as the primary source of help.
- [Course and conference](#) information, including extensive reference material.
- [General information](#) about the project.
- [Package developer](#) resources, including guidelines for creating and submitting new packages.

### Exercise 3

Scavenger hunt. Spend five minutes tracking down the following information.

- a. From the *Bioconductor* web site, instructions for installing or updating *Bioconductor* packages.




Figure 3.1: *Bioconductor* Google analytics, 1-month access, 10 December 2012. Left: access by country. Right: daily access in 2011 (orange) and 2012 (blue).

- b. A list of all packages in the current release of *Bioconductor*.
- c. The URL of the *Bioconductor* mailing list subscription page.

**Solution:** Possible solutions from the *Bioconductor* web site are, e.g., <http://bioconductor.org/install/> (installation instructions), <http://bioconductor.org/packages/release/bioc/> (current software packages), <http://bioconductor.org/help/mailling-list/> (mailing lists).

## 3.1 High-throughput sequence analysis

Recent technological developments introduce high-throughput sequencing approaches. A variety of experimental protocols and analysis work flows address gene expression, regulation, and encoding of genetic variants. Experimental protocols produce a large number (tens of millions per sample) of short (e.g., 35-150, single or paired-end) nucleotide sequences. These are aligned to a reference or other genome. Analysis work flows use the alignments to infer levels of gene expression (RNA-seq), binding of regulatory elements to genomic locations (ChIP-seq), or prevalence of structural variants (e.g., SNPs, short indels, large-scale genomic rearrangements). Sample sizes range from minimal replication (e.g., 2 samples per treatment group) to thousands of individuals.

## 3.2 Statistical programming

Many academic and commercial software products are available; why would one use *R* and *Bioconductor*? One answer is to ask about the demands high-throughput genomic data places on effective computational biology software.

**Effective computational biology software** High-throughput questions make use of large data sets. This applies both to the primary data (microarray expression values, sequenced reads, etc.) and also to the annotations on those data (coordinates of genes and features such as exons or regulatory regions; participation in biological pathways, etc.). Large data sets place demands on our tools that preclude some standard approaches, such as spread sheets. Likewise, intricate relationships between data and annotation, and the diversity of research questions, require flexibility typical of a programming language rather than a narrowly-enabled graphical user interface.

Analysis of high-throughput data is necessarily statistical. The volume of data requires that it be appropriately summarized before any sort of comprehension is possible. The data are produced by advanced technologies, and these introduce artifacts (e.g., probe-specific bias in microarrays; sequence or base calling bias in RNA-seq experiments) that need to be accommodated to avoid incorrect or inefficient inference. Data sets typically derive from designed experiments, requiring a statistical approach both to account for the design and to correctly address the large number of observed values (e.g., gene expression or sequence tag counts) and small number of samples accessible in typical experiments.

Research needs to be reproducible. Reproducibility is both an ideal of the scientific method, and a pragmatic requirement. The latter comes from the long-term and multi-participant nature of contemporary science. An analysis will be performed for the initial experiment, revisited again during manuscript preparation, and revisited during reviews or in determining next steps. Likewise, analysis typically involve a team of individuals with diverse domains of expertise. Effective collaborations result when it is easy to reproduce, perhaps with minor modifications, an existing result, and when sophisticated statistical or bioinformatic analysis can be effectively conveyed to other group members.

Science moves very quickly. This is driven by the novel questions that are the hallmark of discovery, and by technological innovation and accessibility. Rapidity of scientific development places significant burdens on software, which must also move quickly. Effective software cannot be too polished, because that requires that the correct analyses are ‘known’ and that significant resources of time and money have been invested in developing the software; this implies software that is tracking the trailing edge of innovation. On the other

hand, leading-edge software cannot be too idiosyncratic; it must be usable by a wider audience than the creator of the software, and fit in with other software relevant to the analysis.

Effective software must be accessible. Affordability is one aspect of accessibility. Another is transparent implementation, where the novel software is sufficiently documented and source code accessible enough for the assumptions, approaches, practical implementation decisions, and inevitable coding errors to be assessed by other skilled practitioners. A final aspect of affordability is that the software is actually usable. This is achieved through adequate documentation, support forums, and training opportunities.

***Bioconductor* as effective computational biology software** What features of *R* and *Bioconductor* contribute to its effectiveness as a software tool?

*Bioconductor* is well suited to handle extensive data and annotation. *Bioconductor* ‘classes’ represent high-throughput data and their annotation in an integrated way. *Bioconductor* methods use advanced programming techniques or *R* resources (such as transparent data base or network access) to minimize memory requirements and integrate with diverse resources. Classes and methods coordinate complicated data sets with extensive annotation. Nonetheless, the basic model for object manipulation in *R* involves vectorized in-memory representations. For this reason, particular programming paradigms (e.g., block processing of data streams; explicit parallelism) or hardware resources (e.g., large-memory computers) are sometimes required when dealing with extensive data.

*R* is ideally suited to addressing the statistical challenges of high-throughput data. Three examples include the development of the ‘RMA’ and other normalization algorithm for microarray pre-processing, use of moderated *t*-statistics for assessing microarray differential expression, and development of negative binomial approaches to estimating dispersion read counts necessary for appropriate analysis of RNAseq designed experiments.

Many of the ‘old school’ aspects of *R* and *Bioconductor* facilitate reproducible research. An analysis is often represented as a text-based script. Reproducing the analysis involves re-running the script; adjusting how the analysis is performed involves simple text-editing tasks. Beyond this, *R* has the notion of a ‘vignette’, which represents an analysis as a  $\text{\LaTeX}$  document with embedded *R* commands. The *R* commands are evaluated when the document is built, thus reproducing the analysis. The use of  $\text{\LaTeX}$  means that the symbolic manipulations in the script are augmented with textual explanations and justifications for the approach taken; these include graphical and tabular summaries at appropriate places in the analysis. *R* includes facilities for reporting the exact version of *R* and associated packages used in an analysis so that, if needed, discrepancies between software versions can be tracked down and their importance evaluated. While users often think of *R* packages as providing new functionality, packages are also used to enhance reproducibility by encapsulating a single analysis. The package can contain data sets, vignette(s) describing the analysis, *R* functions that might have been written, scripts for key data processing stages, and documentation (via standard *R* help mechanisms) of what the functions, data, and packages are about.

The *Bioconductor* project adopts practices that facilitate reproducibility. Versions of *R* and *Bioconductor* are released twice each year. Each *Bioconductor* release is the result of development, in a separate branch, during the previous six months. The release is built daily against the corresponding version of *R* on Linux, Mac, and Windows platforms, with an extensive suite of tests performed. The `biocLite` function ensures that each release of *R* uses the corresponding *Bioconductor* packages. The user thus has access to stable and tested package versions. *R* and *Bioconductor* are effective tools for reproducible research.

*R* and *Bioconductor* exist on the leading portion of the software life cycle. Contributors are primarily from academic institutions, and are directly involved in novel research activities. New developments are made available in a familiar format, i.e., the *R* language, packaging, and build systems. The rich set of facilities in *R* (e.g., for advanced statistical analysis or visualization) and the extensive resources in *Bioconductor* (e.g., for annotation using third-party data such as Biomart or UCSC genome browser tracks) mean that innovations can be directly incorporated into existing work flows. The ‘development’ branches of *R* and *Bioconductor* provide an environment where contributors can explore new approaches without alienating their user base.

*R* and *Bioconductor* also fair well in terms of accessibility. The software is freely available. The source

Table 3.1: Selected *Bioconductor* packages for high-throughput sequence analysis.

Concept	Packages
Data representation	<i>IRanges</i> , <i>GenomicRanges</i> , <i>GenomicFeatures</i> , <i>Biostrings</i> , <i>BSgenome</i> , <i>girafe</i> .
Input / output	<i>ShortRead</i> (fastq), <i>Rsamtools</i> (bam), <i>rtracklayer</i> (gff, wig, bed), <i>VariantAnnotation</i> (vcf), <i>R453Plus1Toolbox</i> (454).
Annotation	<i>GenomicFeatures</i> , <i>ChIPpeakAnno</i> , <i>VariantAnnotation</i> .
Alignment	<i>gmapR</i> , <i>Rsubread</i> , <i>Biostrings</i> .
Visualization	<i>ggbio</i> , <i>Gviz</i> .
Quality assessment	<i>qRQC</i> , <i>seqbias</i> , <i>ReQON</i> , <i>htSeqTools</i> , <i>TEQC</i> , <i>Rolexa</i> , <i>ShortRead</i> .
RNA-seq	<i>BitSeq</i> , <i>cqn</i> , <i>cummeRbund</i> , <i>DESeq</i> , <i>DEXSeq</i> , <i>EDASeq</i> , <i>edgeR</i> , <i>gage</i> , <i>goseq</i> , <i>iASeq</i> , <i>tweeDEseq</i> .
ChIP-seq, etc.	<i>BayesPeak</i> , <i>baySeq</i> , <i>ChIPpeakAnno</i> , <i>chipseq</i> , <i>ChIPseqR</i> , <i>ChIPsim</i> , <i>CSAR</i> , <i>DiffBind</i> , <i>MEDIPS</i> , <i>mosaics</i> , <i>NarrowPeaks</i> , <i>nucleR</i> , <i>PICS</i> , <i>PING</i> , <i>REDseq</i> , <i>Repitools</i> , <i>TSSi</i> .
Variants	<i>VariantAnnotation</i> , <i>VariantTools</i> , <i>gmapR</i>
SNPs	<i>snpStats</i> , <i>GWASTools</i> , <i>hapFabia</i> , <i>GGtools</i>
Copy number	<i>cn.mops</i> , <i>genoset</i> , <i>CNAnorm</i> , <i>exomeCopy</i> , <i>segmentSeq</i> .
Motifs	<i>MotifDb</i> , <i>BCRANK</i> , <i>cosmo</i> , <i>cosmoGUI</i> , <i>MotIV</i> , <i>seqLogo</i> , <i>rGADEM</i> .
3C, etc.	<i>HiTC</i> , <i>r3Cseq</i> .
Microbiome	<i>phyloseq</i> , <i>DirichletMultinomial</i> , <i>clstutils</i> , <i>manta</i> , <i>mcaGUI</i> .
Work flows	<i>QuasR</i> , <i>easyRNASeq</i> , <i>ArrayExpressHTS</i> , <i>Genominator</i> , <i>oneChannel-GUI</i> , <i>rnaSeqMap</i> .
Database	<i>SRADB</i> .

code is easily and fully accessible for critical evaluation. The *R* packaging and check system requires that all functions are documented. *Bioconductor* requires that each package contain vignettes to illustrate the use of the software. There are very active *R* and *Bioconductor* mailing lists for immediate support, and regular training and conference activities for professional development.

### 3.3 *Bioconductor* packages for high-throughput sequence analysis

Table 3.1 enumerates many of the packages available for sequence analysis. The table includes packages for representing sequence-related data (e.g., *GenomicRanges*, *Biostrings*), as well as domain-specific analysis such as RNA-seq (e.g., *edgeR*, *DEXSeq*), ChIP-seq (e.g., *ChIPpeakAnno*, *DiffBind*), variants (e.g., *VariantAnnotation*, *VariantTools*), and SNPs and copy number variation (e.g., *genoset*, *ggtools*).

### 3.4 S4 Classes and methods

*Bioconductor* makes extensive use of ‘S4’ classes. Essential operations are sketched in Table 3.2. *Bioconductor* has tried to develop and use ‘best practices’ for use of S4 classes. Usually instances are created by a call to a *constructor*, such as *GRanges* (an object representing genomic ranges, with information on sequence, strand, start, and end coordinate of each range), or are returned by a function call that makes the object ‘behind the scenes’ (e.g., *readFastq*). Objects can have complicated structure, but users are not expected to have to concern themselves with the internal representation, just as the details of the S3 object returned by the *lm* function are not of direct concern. Instead, one might query the object to retrieve information; functions providing this functionality are sometimes called *accessors*, e.g., *seqnames*; the data that is returned by the accessor may involve some calculation, e.g., querying a data base, that the user can remain blissfully unaware of. It can be important to appreciate that objects can be related to one another, in particular inheriting

Table 3.2: Using S4 classes and methods.

Best practices	
<code>gr &lt;- GRanges()</code>	‘Constructor’; create an instance of the <i>GRanges</i> class
<code>seqnames(gr)</code>	‘Accessor’, extract information from an instance
<code>countOverlaps(gr1, gr2)</code>	A <i>method</i> implementing a <i>generic</i> with useful functionality
Older packages	
<code>s &lt;- new("MutliSet")</code>	A constructor
<code>s@annotation</code>	A ‘slot’ accessor
Help	
<code>class(gr)</code>	Discover class of instance
<code>getClass(gr)</code>	Display class structure, e.g., inheritance
<code>showMethods(findOverlaps)</code>	Classes for which methods of <code>findOverlaps</code> are implemented
<code>showMethods(class="GRanges", where=search())</code>	Generics with methods implemented for the <i>GRanges</i> class, limited to currently loaded packages.
<code>class?GRanges</code>	Documentation for the <i>GRanges</i> class.
<code>method?"findOverlaps,GRanges,GRanges"</code>	Documentation for the <code>findOverlaps</code> method when the two arguments are both <i>GRanges</i> instances.
<code>selectMethod(findOverlaps, c("GRanges", "GRanges"))</code>	View source code for the method, including method ‘dispatch’

parts of its internal structure and external behavior from other classes. For instance, the *GRanges* class inherits structure and behavior from the *GenomicRanges* class. The details of structural inheritance should not be important to the user, but the fact that once class inherits from another can be useful information to know.

One often calls a function in which one or more objects are arguments, e.g., `countOverlaps` can take two *GRanges* instances. The role of the function is to transform inputs into outputs. In the case of `countOverlaps` the transformation is to summarize the number of ranges in the second argument (the `subject` function argument) overlap with ranges in the first argument (the `query` argument). This establishes a kind of contract, e.g., the return value of `countOverlaps` should be a non-negative integer vector, with as many elements as there are ranges in the `query` argument, and with a one-to-one correspondence between elements in the `query` input argument and the output. Having established such a contract, it can be convenient to write variations of `countOverlaps` that fulfill the contract but for different objects, e.g., when the arguments are instances of class *IRanges*, which do not have information about chromosomal sequence or strand. To indicate that the same contract is being fulfilled, and perhaps to simplify software development, one typically makes `countOverlaps` a *generic* function, and implements methods for different types of arguments.

Attending course and reading vignettes pages are obviously an excellent way to get an initial orientation about classes and methods that are available. It can be very helpful, as one becomes more proficient, to use the interactive help system to discover what can be done with the objects one has or the functions one knows about.

The `showMethods` function is a key entry point into discovery of available methods, e.g., `showMethods("countOverlaps")` to show methods defined on the `countOverlaps` generic, or `showMethods(class="GRanges", where=search())` to discover methods available to transform *GRanges* instances. The definition of a method can be retrieved as

```
> selectMethod(countOverlaps, c("GRanges", "GRanges"))
```

## Exercise 4

Load the *GenomicRanges* package.

- Use `getClass` to discover the class structure of *GRanges*, paying particular attention to inheritance relationships summarized in the “Extends:” section of the display.
- Use `showMethods` to see what methods are defined for the `countOverlaps` function.
- There are many methods defined for `countOverlaps`, but none are listed for the *GRanges*, *GRanges* combination of arguments. Yet `countOverlaps` does work when provided with two *GRanges* arguments. Why is that?

**Solution:** Here we load the package and ask about class structure.

```
> library(GenomicRanges)
> getClass("GRanges")
```

```
Class "GRanges" [package "GenomicRanges"]
```

Slots:

Name:	seqnames	ranges	strand	elementMetadata
Class:	Rle	IRanges	Rle	DataFrame

Name:	seqinfo	metadata
Class:	Seqinfo	list

Extends:

```
Class "GenomicRanges", directly
Class "Vector", by class "GenomicRanges", distance 2
Class "GenomicRangesORmissing", by class "GenomicRanges", distance 2
Class "Annotated", by class "GenomicRanges", distance 3
Class "GenomicRangesORGRangesList", by class "GenomicRanges", distance 2
```

*GRanges* extends several classes, including *GenomicRanges*. Method defined on the `countOverlaps` generic can be discovered with

```
> showMethods("countOverlaps")
```

```
Function: countOverlaps (package IRanges)
query="ANY", subject="missing"
query="ANY", subject="Vector"
query="GappedAlignmentPairs", subject="Vector"
query="GappedAlignments", subject="GappedAlignments"
query="GappedAlignments", subject="Vector"
query="GenomicRanges", subject="GenomicRanges"
query="GenomicRanges", subject="Vector"
query="GRangesList", subject="GRangesList"
query="GRangesList", subject="Vector"
query="RangedData", subject="RangedData"
query="RangedData", subject="RangesList"
query="RangesList", subject="RangedData"
query="RangesList", subject="RangesList"
query="SummarizedExperiment", subject="SummarizedExperiment"
query="SummarizedExperiment", subject="Vector"
query="Vector", subject="GappedAlignments"
```

```

query="Vector", subject="GenomicRanges"
query="Vector", subject="GRangesList"
query="Vector", subject="SummarizedExperiment"
query="Vector", subject="ViewsList"
query="ViewsList", subject="Vector"
query="ViewsList", subject="ViewsList"

```

(quotation marks are, in this case, optional). Note that there is no method defined for the *GRanges*, *GRanges* combination of arguments. Yet `countOverlaps` *does* work... (skim over the details of the *se* objects; we are using a constructor to make genomic ranges on plus strand of chromosome 1; there are two ranges in `gr0`, and one in `gr1`).

```

> gr0 <- GRanges("chr1", IRanges(start=c(10, 20), width = 5), "+")
> gr1 <- GRanges("chr1", IRanges(start=12, end=18), "+")
> countOverlaps(gr0, gr1)

[1] 1 0

```

`gr1` overlaps the first range of `gr0`, but not the second, and we end up with a vector of counts `c(1, 0)`. The reason that this ‘works’ is because of inheritance – *GRanges* extends *GenomicRanges*, and we end up selecting the inherited method *countOverlaps, GenomicRanges, GenomicRanges-method*.

## 3.5 Help!

**S4 classes, generics, and methods** To illustrate how help work with S4 classes and generics, consider the *DNASTringSet* class complement generic in the *Biostrings* package:

```

> library(Biostrings)
> showMethods(complement)

Function: complement (package Biostrings)
x="DNASTring"
x="DNASTringSet"
x="MaskedDNASTring"
x="MaskedRNAString"
x="RNAString"
x="RNAStringSet"
x="XStringViews"

```

(Most) methods defined on the *DNASTringSet* class of *Biostrings* and available on the current search path can be found with

```

> showMethods(class="DNASTringSet", where=search())

```

Obtaining help on S4 classes and methods requires syntax such as

```

> class ? DNASTringSet
> method ? "complement,DNASTringSet"

```

The specification of method and class in the latter must not contain a space after the comma.



# Chapter 4

## Sequencing

### 4.1 Technologies

The most common ‘second generation’ technologies readily available to labs are

- Illumina single- and paired-end reads. Short ( $\approx 100$  per end) and very numerous. Flow cell, lane, bar-code.
- Roche 454. 100’s of nucleotides, 100,000’s of reads.
- Life Technologies SOLiD. Unique ‘color space’ model.
- Complete Genomics. Whole genome sequence / variants / etc as a service; end user gets derived results.

Figure 4.1 illustrates Illumina and 454 sequencing. *Bioconductor* has good support for Illumina and Roche 454 sequencing products, and for derived data such as aligned reads or called variants; use of SOLiD color space reads typically requires conversion to FASTQ files that undermine the benefit of the color space model.

All second-generation technologies rely on PCR and other techniques to generate reads from samples that represent aggregations of many DNA molecules. ‘Third-generation’ technologies shift to single-molecule sequencing, with relevant players including Pacific Biosciences and IonTorrent. This data is not widely available, and will not be discussed further.

The most common data in *Bioconductor* work flows is from Illumina sequencers. Reads are either single-end or paired-end. Single-end reads represent 30– $\approx 100$  nucleotides sequenced from DNA that has been sheared into  $\approx 300$  nucleotide fragments. Paired-end reads represent 30– $\approx 100$  nucleotide reads that are paired, and from both ends of the  $\approx 300$  fragment.

Sequence data can be derived from a tremendous diversity of experiments. Some of the most common include:

**RNA-seq** Sequencing of reverse-complemented mRNA from the entire expressed transcriptome, typically.

Used for *differential expression* studies like micro-arrays, or for *novel transcript discovery*.

**DNA-seq** Sequencing of whole or targeted (e.g., exome) genomic DNA. Common goals include *SNP* detection, *indel* and other structural polymorphisms, and *CNV* (copy number variation). DNA-seq is also used for *de novo* assembly, but *de novo* assembly is not an area where *Bioconductor* contributes.

**ChIP-seq** ChIP (chromatin immuno-precipitation) is used to enrich genomic DNA for regulatory elements, followed by sequencing and mapping of the enriched DNA to a reference genome. The initial statistical challenge is to identify regions where the mapped reads are enriched relative to a sample that did not undergo ChIP[18]; a subsequent task is to identify differential binding across a designed experiment, e.g., [21].

**Metagenomics** Sequencing generates sequences from samples containing multiple species, typically microbial communities sampled from niches such as the human oral cavity. Goals include inference of *species composition* (when sequencing typically targets phylogenetically informative genes such as 16S) or *metabolic contribution*.




Figure 4.1: High-throughput sequencing. Left: Illumina bridge PCR [2]; mis-call errors. Right: Roche 454 [22]; homopolymer errors.

## 4.2 Data

### 4.2.1 A running example: the *pasilla* data set

As a running example, we use the *pasilla* data set, derived from [3]. The authors investigate conservation of RNA regulation between *D. melanogaster* and mammals. Part of their study used RNAi and RNA-seq to identify exons regulated by Pasilla (*ps*), the *D. melanogaster* ortholog of mammalian NOVA1 and NOVA2. Briefly, their experiment compared gene expression as measured by RNAseq in S2-DRSC cells cultured with, or without, a 444bp dsRNA fragment corresponding to the *ps* mRNA sequence. Their assessment investigated differential exon use, but our worked example will focus on gene-level differences. For several examples we look at a subset of the *ps* data, corresponding to reads obtained from lanes of their RNA-seq experiment, and to the same reads aligned to a *D. melanogaster* reference genome. Reads were obtained from GEO and the Short Read Archive (SRA), and were aligned to the *D. melanogaster* reference genome *dm3* as described in the *pasilla* experiment data package.

### 4.2.2 Work flows

At a very high level, one can envision a work flow that starts with a challenging biological question (how does *ps* influence gene and transcript regulation?). The biological question is framed in terms of wet-lab protocols coupled with an appropriate and all-important experimental design. There are several well-known statistical at this stage, common to any experimental data What treatments are going to be applied? How many replicates will there be of each? Is there likely to be sufficient power to answer the biologically relevant question? Reality is also important at this stage, as evidenced in the *pasilla* data where, as we will see, samples were collected using different methods (single versus paired end reads) over a time when there were rapid technological changes. Such reality often introduces confounding factors that require appropriate statistical treatment in subsequent analysis.

The work flow proceeds with data generation, involving both a wet-lab (sample preparation) component and actual sequencing. It is essential to acknowledge the biases and artifacts that are introduced at each of these stages. Sample preparation involves non-trivial amounts of time and effort. Large studies are likely to have batch effects (e.g., because work was done by different lab members, or different batches of reagent). Samples might have been prepared in ways that are likely to influence down-stream analysis, e.g., using a protocol involving PCR and hence introducing opportunities for sample-specific bias. DNA isolation protocols may introduce many artifacts, e.g., non-uniform representation of reads across the length

Table 4.1: Common file types (e.g., <http://genome.ucsc.edu/FAQ/FAQformat.html>) and *Bioconductor* packages used for input.

Description	File	Package
Unaligned sequences: identifier, sequence, and encoded quality score tuples	FASTQ	<i>ShortRead</i>
Aligned sequences: identifier, sequence, reference sequence name, strand position, cigar and additional tags	BAM	<i>Rsamtools</i>
Gene annotations: reference sequence name, data source, feature type, start and end positions, strand, etc.	GFF, GTF	<i>rtracklayer</i>
Range-based annotation: reference sequence name, start, end coordinates.	BED	<i>rtracklayer</i>
‘Continuous’ single-nucleotide annotation.	WIG, bigWig	<i>rtracklayer</i>
Called single nucleotide, indel, copy number, and structural variants, often compressed and indexed (with <i>Rsamtools</i> <code>bgzip</code> , <code>indexTabix</code> )	VCF	<i>VariantAnnotation</i>

of expressed genes in RNA-seq. The sequencing reaction itself is far from bias-free, with known artifacts of called base frequency, cycle-dependent accuracy and bias, non-uniform coverage, etc. At a minimum, the research needs to be aware of the opportunities for bias that can be introduced during sample preparation and sequencing.

The informatics component of work flows becomes increasingly important during and after sequence generation. The sequencer is often treated as a ‘black box’, producing short reads consisting of 10’s to 100’s of nucleotides and with associated quality scores. Usually, the chemistry and informatics processing pipeline are sufficiently well documented that one can arrive at an understanding of biases and quality issues that might be involved; such an understanding is likely to be particularly important when embarking on questions or using protocols that are at the fringe of standard practice (where, after all, the excitement is).

The first real data seen by users are *fastq* files (Table 4.1). These files are often simple text files consisting of many millions of records, and are described in greater detail in Section 5.2. The center performing the sequencing typically vets results for quality, but these quality measures are really about the performance of their machines. It is very important to assess quality with respect to the experiment being undertaken – Are the numbers of reads consistent across samples? Is the GC content and other observable aspects of the reads consistent with expectation? Are there anomalies in the sequence results that reflect primers or other reagents used during sample preparation? Are well-known artifacts of the protocol used evident in the reads in hand?

The next step in many work flows involves alignment of reads to a reference genome. There are many aligners available, including *BWA* [12], *Bowtie* / *Bowtie2* [11], and *GSNAP*; merits of these are discussed in the literature. *Bioconductor* packages ‘wrapping’ these tools are increasingly common (e.g., *Rbowtie*, *gmapR*; *cummeRbund* for parsing output of the *cufflinks* transcript discovery pathway). There are also alignment algorithms implemented in *Bioconductor* (e.g., `matchPDict` in the *Biostrings* package, and the *Rsubread* package); `matchPDict` is particularly useful for flexible alignment of moderately sized subsets of data. Most main-stream aligners produce output in ‘SAM’ or ‘BAM’ (binary alignment) format. BAM files are the primary starting point for many analyses, and their manipulation and use in *Bioconductor* is introduced in Section 6.2.

Common analyses often use well-established third-party tools for initial stages of the analysis; some of these have *Bioconductor* counterparts that are particularly useful when the question under investigation does not meet the assumptions of other facilities. Some common work flows (a more comprehensive list is available on the SeqAnswers wiki<sup>1</sup>) include:

**ChIP-seq** ChIP-seq experiments typically use DNA sequencing to identify regions of genomic DNA enriched

<sup>1</sup><http://seqanswers.com/wiki/RNA-Seq>

in prepared samples relative to controls. A central task is thus to identify peaks, with common tools including *MACS* and *PeakRanger*.

**RNA-seq** In addition to the aligners mentioned above, RNA-seq for differential expression might use the *HTSeq*<sup>2</sup> python tools for counting reads within regions of interest (e.g., known genes) or a pipeline such as the *bowtie* (basic alignment) / *tophat* (splice junction mapper) / *cufflinks* (estimated isoform abundance) (e.g., <sup>3</sup>) or *RSEM*<sup>4</sup> suite of tools for estimating transcript abundance.

**DNA-seq** especially variant calling can be facilitated by software such as the *GATK*<sup>5</sup> toolkit.

There are many *R* packages that replace or augment the analyses outlined above, as summarized in Table 3.1.

Programs such as those outlined the previous paragraph often rely on information about gene or other structure as input, or produce information about chromosomal locations of interesting features. The GTF and BED file formats are common representations of this information. Representing these files as *R* data structures is often facilitated by the *rtracklayer* package. We explore these files in Chapter 12. Variants are very commonly represented in VCF (Variant Call Format) files; these are explored in Chapter 10.

---

<sup>2</sup><http://www-huber.embl.de/users/anders/HTSeq/doc/overview.html>

<sup>3</sup><http://bowtie-bio.sourceforge.net/index.shtml>

<sup>4</sup><http://deweylab.biostat.wisc.edu/rsem/>

<sup>5</sup><http://www.broadinstitute.org/gatk/>

## Chapter 5

# Strings and Reads

### 5.1 DNA (and other) Strings with the *Biostrings* package

The *Biostrings* package provides tools for working with sequences. The essential data structures are *DNASTring* and *DNASTringSet*, for working with one or multiple DNA sequences. The *Biostrings* package contains additional classes for representing amino acid and general biological strings. The *BSgenome* and related packages (e.g., *BSgenome.Dmelanogaster.UCSC.dm3*) are used to represent whole-genome sequences. Table 5.2 summarizes common operations; The following exercise explores these packages.

#### Exercise 5

The objective of this exercise is to calculate the GC content of the exons of a single gene. We jump into the middle of some of the data structures common in Bioconductor; these are introduced more thoroughly in later exercises..

Load the *BSgenome.Dmelanogaster.UCSC.dm3* data package, containing the UCSC representation of *D. melanogaster* genome assembly dm3. Discover the content of the package by evaluating *Dmelanogaster*.

Load the *SequenceAnalysisData* package, and evaluate the command `data(ex)` to load an example of a *GRangesList* object. the *GRangesList* represents coordinates of exons in the *D. melanogaster* genome, grouped by gene.

Look at `ex[1]`. These are the genomic coordinates of the first gene in the *ex* object. Load the *D. melanogaster* chromosome that this gene is on by subsetting the *Dmelanogaster* object.

Use *Views* to create views on to the chromosome that span the start and end coordinates of all exons in the first gene; the start and end coordinates are accessed with `start(ex[[1]])` and similar.

Develop a function *gcFunction* to calculate GC content. Use this to calculate the GC content in each of the exons.

**Solution:** Here we load the *D. melanogaster* genome, select a single chromosome, and create *Views* that reflect the ranges of the *FBgn0002183*.

Table 5.1: Selected Bioconductor packages for representing strings and reads.

Package	Description
<i>Biostrings</i>	Classes (e.g., <i>DNASTringSet</i> ) and methods (e.g., <code>alphabetFrequency</code> , <code>pairwiseAlignment</code> ) for representing and manipulating DNA and other biological sequences.
<i>BSgenome</i>	Representation and manipulation of large (e.g., whole-genome) sequences.
<i>ShortRead</i>	I/O and manipulation of FASTQ files.

Table 5.2: Operations on strings in the *Biostrings* package.

	Function	Description
Access	<code>length, names</code>	Number and names of sequences
	<code>[, head, tail, rev</code>	Subset, first, last, or reverse sequences
	<code>c</code>	Concatenate two or more objects
	<code>width, nchar</code>	Number of letters of each sequence
	<code>Views</code>	Light-weight sub-sequences of a sequence
Compare	<code>==, !=, match, %in%</code>	Element-wise comparison
	<code>duplicated, unique</code>	Analog to <code>duplicated</code> and <code>unique</code> on character vectors
	<code>sort, order</code>	Locale-independent sort, order
	<code>split, relist</code>	Split or relist objects to, e.g., <i>DNAStringSetList</i>
Edit	<code>subseq, subseq&lt;-</code>	Extract or replace sub-sequences in a set of sequences
	<code>reverse, complement</code>	Reverse, complement, or reverse-complement DNA
	<code>reverseComplement</code>	
	<code>translate</code>	Translate DNA to Amino Acid sequences
	<code>chartr</code>	Translate between letters
	<code>replaceLetterAt</code>	Replace letters at a set of positions by new letters
	<code>trimLRPatterns</code>	Trim or find flanking patterns
	<code>alphabetFrequency</code>	Tabulate letter occurrence
Count	<code>letterFrequency</code>	
	<code>letterFrequencyInSlidingView</code>	
	<code>consensusMatrix</code>	Nucleotide $\times$ position summary of letter counts
	<code>dinucleotideFrequency</code>	2-mer, 3-mer, and k-mer counting
	<code>trinucleotideFrequency</code>	
	<code>oligonucleotideFrequency</code>	
	<code>nucleotideFrequencyAt</code>	Nucleotide counts at fixed sequence positions
	<code>matchPattern, countPattern</code>	Short patterns in one or many ( <i>v*</i> ) sequences
	<code>vmatchPattern, vcountPattern</code>	
	<code>matchPDict, countPDict</code>	Short patterns in one or many ( <i>v*</i> ) sequences (mismatch only)
Match	<code>whichPDict, vcountPDict</code>	
	<code>vwhichPDict</code>	
	<code>pairwiseAlignment</code>	Needleman-Wunsch, Smith-Waterman, etc. pairwise alignment
	<code>matchPWM, countPWM</code>	Occurrences of a position weight matrix
	<code>matchProbePair</code>	Find left or right flanking patterns
	<code>findPalindromes</code>	Palindromic regions in a sequence. Also
	<code>findComplementedPalindromes</code>	
	<code>stringDist</code>	Levenshtein, Hamming, or pairwise alignment scores
	<code>readDNAStringSet</code>	FASTA (or sequence only from FASTQ). Also
	<code>readBStringSet, readRNAStringSet, readAAStringSet</code>	
I/O	<code>writeXStringSet</code>	
	<code>writePairwiseAlignments</code>	Write <code>pairwiseAlignment</code> as “pair” format
	<code>readDNAMultipleAlignment</code>	Multiple alignments (FASTA, “stockholm”, or “clustal”). Also
	<code>readRNAMultipleAlignment, readAAMultipleAlignment</code>	
	<code>write.phylip</code>	

```

> library(BSgenome.Dmelanogaster.UCSC.dm3)
> Dmelanogaster

Fly genome
|
| organism: Drosophila melanogaster (Fly)
| provider: UCSC
| provider version: dm3
| release date: Apr. 2006
| release name: BDGP Release 5
|
| single sequences (see '?seqnames'):
| chr2L      chr2R      chr3L      chr3R      chr4      chrX      chrU
| chrM      chr2LHet   chr2RHet   chr3LHet   chr3RHet   chrXHet   chrYHet
| chrUextra
|
| multiple sequences (see '?mseqnames'):
| upstream1000 upstream2000 upstream5000
|
| (use the '$' or '[' operator to access a given sequence)

> library(SequenceAnalysisData)
> data(ex)
> ex[1]

GRangesList of length 1:
$FBgn0002183
GRanges with 9 ranges and 0 metadata columns:
      seqnames      ranges strand
      <Rle>         <IRanges> <Rle>
[1] chr3L [1871574, 1871917] -
[2] chr3L [1872354, 1872470] -
[3] chr3L [1872582, 1872735] -
[4] chr3L [1872800, 1873062] -
[5] chr3L [1873117, 1873983] -
[6] chr3L [1874041, 1875218] -
[7] chr3L [1875287, 1875586] -
[8] chr3L [1875652, 1875915] -
[9] chr3L [1876110, 1876336] -

---
seqlengths:
      chr2L chr2LHet      chr2R chr2RHet ... chrXHet chrYHet      chrM
23011544  368872 21146708  3288761 ...  204112  347038 19517

> nm <- "chr3L"
> chr <- Dmelanogaster[[nm]]
> v <- Views(chr, start=start(ex[[1]]), end=end(ex[[1]]))

Here is the gcFunction helper function to calculate GC content:

> gcFunction <-
+   function(x)

```

```
+ {
+   alf <- alphabetFrequency(x, as.prob=TRUE)
+   rowSums(alf[,c("G", "C")])
+ }
```

The `gcFunction` is really straight-forward: it invokes the function `alphabetFrequency` from the [Biostrings](#) package. This returns a simple matrix of exon  $\times$  nucleotide probabilities. The row sums of the G and C columns of this matrix are the GC contents of each exon.

The subject GC content is

```
> subjectGC <- gcFunction(v)
```

## 5.2 Reads and the [ShortRead](#) package

**Short read formats** The Illumina GAII and HiSeq technologies generate sequences by measuring incorporation of florescent nucleotides over successive PCR cycles. These sequencers produce output in a variety of formats, but *FASTQ* is ubiquitous. Each read is represented by a record of four components:

```
@SRR031724.1 HWI-EAS299_4_30M2BAAXX:5:1:1513:1024 length=37
GTTTGTCCAAGTTCTGCTAGCTGAATCCTGGGGCGC
+SRR031724.1 HWI-EAS299_4_30M2BAAXX:5:1:1513:1024 length=37
IIIIIIIIIIIIIIIIIIIIIIIIIIII+HIIII<IE
```

The first and third lines (beginning with @ and + respectively) are unique identifiers. The identifier produced by the sequencer typically includes a [machine id](#) followed by colon-separated information on the [lane, tile, x, and y coordinate of the read](#). The example illustrated here also includes the [SRA accession number](#), added when the data was submitted to the archive. The machine identifier could potentially be used to extract information about batch effects. The spatial coordinates (lane, tile, x, y) are often used to identify optical duplicates; spatial coordinates can also be used during quality assessment to identify artifacts of sequencing, e.g., uneven amplification across the flow cell, though these spatial effects are rarely pursued.

The second and fourth lines of the FASTQ record are the nucleotides and qualities of each cycle in the read. This information is given in 5' to 3' orientation as seen by the sequencer. A letter N in the sequence is used to signify bases that the sequencer was not able to call. The fourth line of the FASTQ record encodes the quality (confidence) of the corresponding base call. The quality score is encoded following one of several conventions, with the general notion being that letters later in the visible ASCII alphabet

```
!"#$%&'()*+,-./0123456789:;<=>?@ABCDEFGHIJKLMNO
PQRSTUVWXYZ[\]^_`abcdefghijklmnopqrstuvwxyz{|}~
```

are of higher quality; this is developed further below. Both the sequence and quality scores may span multiple lines.

Technologies other than Illumina use different formats to represent sequences. Roche 454 sequence data is generated by ‘flowing’ labeled nucleotides over samples, with greater intensity corresponding to longer runs of A, C, G, or T. This data is represented as a series of ‘flow grams’ (a kind of run-length encoding of the read) in Standard Flowgram Format (SFF). The *Bioconductor* package [R453Plus1Toolbox](#) has facilities for parsing SFF files, but after quality control steps the data are frequently represented (with some loss of information) as FASTQ. SOLiD technologies produce sequence data using a ‘color space’ model. This data is not easily read in to R, and much of the error-correcting benefit of the color space model is lost when converted to FASTQ; SOLiD sequences are not well-handled by *Bioconductor* packages.



**Short reads in R** FASTQ files can be read in to R using the `readFastq` function from the [ShortRead](#) package. Use this function by providing the path to a FASTQ file. There are sample data files available in the *SequenceAnalysisData* package, each consisting of 1 million reads from a lane of the Pasilla data set.

```
> library(ShortRead)
> bigdata <- system.file("bigdata", package="SequenceAnalysisData")
> fastqDir <- file.path(bigdata, "fastq")
> fastqFiles <- dir(fastqDir, full=TRUE)
> fq <- readFastq(fastqFiles[1])
> fq
```

```
class: ShortReadQ
length: 1000000 reads; width: 37 cycles
```

The data are represented as an object of class *ShortReadQ*.

```
> head(sread(fq), 3)

A DNASTringSet instance of length 3
width seq
[1] 37 GTTTTGTCCAAGTTCTGGTAGCTGAATCCTGGGGCGC
[2] 37 GTTGTGCGATTCCTTACTCTCATTCGGGAATTCTGTT
[3] 37 GAATTTTTTGAGAGCGAAATGATAGCCGATGCCCTGA

> head(quality(fq), 3)
```

```
class: FastqQuality
quality:
A BStringSet instance of length 3
width seq
[1] 37 IIIIIIIIIIIIIIIIIIIIIIIIIII+HIIII<IE
[2] 37 IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII
[3] 37 IIIIIIIIIIIIIIIIIIIIIII'IIIIIGBIIII2I+

> head(id(fq), 3)
```

```
A BStringSet instance of length 3
width seq
[1] 58 SRR031724.1 HWI-EAS299_4_30M2BAAXX:5:1:1513:1024 length=37
[2] 57 SRR031724.2 HWI-EAS299_4_30M2BAAXX:5:1:937:1157 length=37
[3] 58 SRR031724.4 HWI-EAS299_4_30M2BAAXX:5:1:1443:1122 length=37
```

The *ShortReadQ* class illustrates *class inheritance*. It extends the *ShortRead* class

```
> getClass("ShortReadQ")

Class "ShortReadQ" [package "ShortRead"]

Slots:

Name:      quality      sread      id
Class: QualityScore DNASTringSet BStringSet

Extends:
Class "ShortRead", directly
Class ".ShortReadBase", by class "ShortRead", distance 2

Known Subclasses: "AlignedRead"
```

Methods defined on *ShortRead* are available for *ShortReadQ*.

```
> showMethods(class="ShortRead", where="package:ShortRead")
```

For instance, the `width` can be used to demonstrate that all reads consist of 37 nucleotides.

```
> table(width(fq))
```

```
37
1000000
```

The `alphabetByCycle` function summarizes use of nucleotides at each cycle in a (equal width) *ShortReadQ* or *DNAStringSet* instance.

```
> abc <- alphabetByCycle(sread(fq))
> abc[1:4, 1:8]
```

	cycle							
alphabet	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]
A	78194	153156	200468	230120	283083	322913	162766	220205
C	439302	265338	362839	251434	203787	220855	253245	287010
G	397671	270342	258739	356003	301640	247090	227811	246684
T	84833	311164	177954	162443	211490	209142	356178	246101

FASTQ files are getting larger. A very common reason for looking at data at this early stage in the processing pipeline is to explore sequence quality. In these circumstances it is often not necessary to parse the entire FASTQ file. Instead create a representative sample

```
> sampler <- FastqSampler(fastqFiles[1], 1000000)
> yield(sampler) # sample of 1000000 reads
```

```
class: ShortReadQ
length: 1000000 reads; width: 37 cycles
```

A second common scenario is to pre-process reads, e.g., trimming low-quality tails, adapter sequences, or artifacts of sample preparation. The *FastqStreamer* class can be used to ‘stream’ over the fastq files in chunks, processing each chunk independently.

*ShortRead* contains facilities for quality assessment of FASTQ files. Here we generate a report from a sample of 1 million reads from each of our files and display it in a web browser

```
> qas0 <- Map(function(fl, nm) {
+   fq <- FastqSampler(fl)
+   qa(yield(fq), nm)
+ }, fastqFiles,
+   sub("_subset.fastq", "", basename(fastqFiles)))
> qas <- do.call(rbind, qas0)
> rpt <- report(qas, dest=tempfile())
> browseURL(rpt)
```

A report from a larger subset of the experiment is available

```
> rpt <- system.file("GSM461176_81_qa_report", "index.html",
+   package="SequenceAnalysisData")
> browseURL(rpt)
```

## Exercise 6

Use the file path `bigdata` and the `file.path` and `dir` functions to locate the fastq file from [3] (the file was obtained as described in the `pasilla` experiment data package).

Input the fastq files using `readFastq` from the [ShortRead](#) package.

Use `alphabetFrequency` to summarize the GC content of all reads (hint: use the `sread` accessor to extract the reads, and the `collapse=TRUE` argument to the `alphabetFrequency` function). Using the helper function `gcFunction` defined elsewhere in this document, draw a histogram of the distribution of GC frequencies across reads.

Use `alphabetByCycle` to summarize the frequency of each nucleotide, at each cycle. Plot the results using `matplot`, from the `graphics` package.

As an advanced exercise, and if on Mac or Linux, use the `parallel` package and `mclapply` to read and summarize the GC content of reads in two files in parallel.

**Solution:** Discovery:

```
> dir(bigdata)

[1] "bam"                  "dm3.ensGene.txdb.sqlite"
[3] "fastq"

> fls <- dir(file.path(bigdata, "fastq"), full=TRUE)
```

Input:

```
> fq <- readFastq(fls[1])
```

GC content:

```
> alf0 <- alphabetFrequency(sread(fq), as.prob=TRUE, collapse=TRUE)
> sum(alf0[c("G", "C")])

[1] 0.5457237
```

A histogram of the GC content of individual reads is obtained with

```
> gc <- gcFunction(sread(fq))
> hist(gc)
```

Alphabet by cycle:

```
> abc <- alphabetByCycle(sread(fq))
> matplot(t(abc[c("A", "C", "G", "T"),]), type="l")
```

Advanced (Mac, Linux only): processing on multiple cores.

```
> library(parallel)
> gc0 <- mclapply(fls, function(fl) {
+   fq <- readFastq(fl)
+   gc <- gcFunction(sread(fq))
+   table(cut(gc, seq(0, 1, .05)))
+ })
> ## simplify list of length 2 to 2-D array
> gc <- simplify2array(gc0)
> matplot(gc, type="s")
```

### Exercise 7

Use `quality` to extract the quality scores of the short reads. Interpret the encoding qualitatively.

Convert the quality scores to a numeric matrix, using `as`. Inspect the numeric matrix (e.g., using `dim`) and understand what it represents.

Use `colMeans` to summarize the average quality score by cycle. Use `plot` to visualize this.

#### Solution:

```
> head(quality(fq))

class: FastqQuality
quality:
  A BStringSet instance of length 6
    width seq
[1] 37 IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII+HIIII<IE
[2] 37 IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII
[3] 37 IIIIIIIIIIIIIIIIIIIIIII'IIIIIGBIIII2I+
[4] 37 IIIIIIIIIIIIIIIIIIIIIIIII,II*E,&4HI++B
[5] 37 IIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIIII&.$
[6] 37 III.IIIIIIIIIIIIIIIIIIIII%IIIE(-EIH<IIII

> qual <- as(quality(fq), "matrix")
> dim(qual)

[1] 1000000      37

> plot(colMeans(qual), type="b")
```

### Exercise 8

As an independent exercise, visit the [qrqc](#) landing page and explore the package vignette. Use the [qrqc](#) package (you may need to install this) to generate base and average quality plots for the data, like those in the report generated by [ShortRead](#).

## Chapter 6

# Ranges and Alignments

Ranges describe both features of interest (e.g., genes, exons, promoters) and reads aligned to the genome. *Bioconductor* has very powerful facilities for working with ranges, some of which are summarized in Table 6.1.

### 6.1 Ranges and the *GenomicRanges* package

Next-generation sequencing data consists of a large number of short reads. These are, typically, aligned to a reference genome. Basic operations are performed on the alignment, asking e.g., how many reads are aligned in a genomic range defined by nucleotide coordinates (e.g., in the exons of a gene), or how many nucleotides from all the aligned reads cover a set of genomic coordinates. How is this type of data, the aligned reads and the reference genome, to be represented in *R* in a way that allows for effective computation?

The *IRanges*, *GenomicRanges*, and *GenomicFeatures* *Bioconductor* packages provide the essential infrastructure for these operations; we start with the *GRanges* class, defined in *GenomicRanges*.

***GRanges*** Instances of *GRanges* are used to specify genomic coordinates. Suppose we wish to represent two *D. melanogaster* genes. The first is located on the positive strand of chromosome 3R, from position 19967117 to 19973212. The second is on the minus strand of the X chromosome, with ‘left-most’ base at 18962306, and right-most base at 18962925. The coordinates are *1-based* (i.e., the first nucleotide on a chromosome is numbered 1, rather than 0), *left-most* (i.e., reads on the minus strand are defined to ‘start’ at the left-most coordinate, rather than the 5’ coordinate), and *closed* (the start and end coordinates are included in the range; a range with identical start and end coordinates has width 1, a 0-width range is represented by the special construct where the end coordinate is one less than the start coordinate).

A complete definition of these genes as *GRanges* is:

Table 6.1: Selected *Bioconductor* packages for representing and manipulating ranges, strings, and other data structures.

Package	Description
<i>IRanges</i>	Defines important classes (e.g., <i>IRanges</i> , <i>Rle</i> ) and methods (e.g., <i>findOverlaps</i> , <i>countOverlaps</i> ) for representing and manipulating ranges of consecutive values. Also introduces <i>DataFrame</i> , <i>SimpleList</i> and other classes tailored to representing very large data.
<i>GenomicRanges</i>	Range-based classes tailored to sequence representation (e.g., <i>GRanges</i> , <i>GRangesList</i> ), with information about strand and sequence name.
<i>GenomicFeatures</i>	Foundation for manipulating data bases of genomic ranges, e.g., representing coordinates and organization of exons and transcripts of known genes.

```
> genes <- GRanges(seqnames=c("3R", "X"),
+                   ranges=IRanges(
+                     start=c(19967117, 18962306),
+                     end=c(19973212, 18962925)),
+                   strand=c("+", "-"),
+                   seqlengths=c(`3R`=27905053L, `X`=22422827L))
```

The components of a *GRanges* object are defined as vectors, e.g., of `seqnames`, much as one would define a *data.frame*. The start and end coordinates are grouped into an *IRanges* instance. The optional `seqlengths` argument specifies the maximum size of each sequence, in this case the lengths of chromosomes 3R and X in the ‘dm2’ build of the *D. melanogaster* genome. This data is displayed as

```
> genes

GRanges with 2 ranges and 0 metadata columns:
      seqnames      ranges strand
      <Rle>        <IRanges> <Rle>
[1]      3R [19967117, 19973212]   +
[2]      X [18962306, 18962925]   -
---
seqlengths:
      3R      X
27905053 22422827
```

For the curious, the gene coordinates and sequence lengths are derived from the [org.Dm.eg.db](#) package for genes with Flybase identifiers FBgn0039155 and FBgn0085359, using the annotation facilities described in Chapter 11.

The *GRanges* class has many useful methods defined on it. Consult the help page

```
> ?GRanges
```

and package vignettes (especially ‘An Introduction to [GenomicRanges](#)’)

```
> vignette(package="GenomicRanges")
```

for a comprehensive introduction. A *GRanges* instance can be subset, with accessors for getting and updating information.

```
> genes[2]

GRanges with 1 range and 0 metadata columns:
      seqnames      ranges strand
      <Rle>        <IRanges> <Rle>
[1]      X [18962306, 18962925]   -
---
seqlengths:
      3R      X
27905053 22422827

> strand(genes)

factor-Rle of length 2 with 2 runs
Lengths: 1 1
Values : + -
Levels(3): + - *
```

```

> width(genes)

[1] 6096 620

> length(genes)

[1] 2

> names(genes) <- c("FBgn0039155", "FBgn0085359")
> genes # now with names

GRanges with 2 ranges and 0 metadata columns:
      seqnames      ranges strand
      <Rle>        <IRanges> <Rle>
FBgn0039155      3R [19967117, 19973212] +
FBgn0085359       X [18962306, 18962925] -
---
seqlengths:
      3R      X
27905053 22422827

```

`strand` returns the strand information in a compact representation called a *run-length encoding*, this is introduced in greater detail below. The ‘names’ could have been specified when the instance was constructed; once named, the *GRanges* instance can be subset by name like a regular vector.

As the *GRanges* function suggests, the *GRanges* class extends the *IRanges* class by adding information about `seqnames`, `strand`, and other information particularly relevant to representing ranges that are on genomes. The *IRanges* class and related data structures (e.g., *RangedData*) are meant as a more general description of ranges defined in an arbitrary space. Many methods implemented on the *GRanges* class are ‘aware’ of the consequences of genomic location, for instance treating ranges on the minus strand differently (reflecting the 5’ orientation imposed by DNA) from ranges on the plus strand.

**Operations on ranges** The *GRanges* class has many useful methods from the *IRanges* class; some of these methods are illustrated here. We use *IRanges* to illustrate these operations to avoid complexities associated with strand and seqnames, but the operations are comparable on *GRanges*. We begin with a simple set of ranges:

```

> ir <- IRanges(start=c(7, 9, 12, 14, 22:24),
+               end=c(15, 11, 12, 18, 26, 27, 28))

```

These and some common operations are illustrated in the upper panel of Figure 6.1 and summarized in Table 6.2.

Methods on ranges can be grouped as follows:

**Intra-range** methods act on each range independently. These include `flank`, `narrow`, `reflect`, `resize`, `restrict`, and `shift`, among others. An illustration is `shift`, which translates each range by the amount specified by the `shift` argument. Positive values shift to the right, negative to the left; `shift` can be a vector, with each element of the vector shifting the corresponding element of the *IRanges* instance. Here we shift all ranges to the right by 5, with the result illustrated in the middle panel of Figure 6.1.

```

> shift(ir, 5)

IRanges of length 7
      start end width
[1]    12  20     9

```




Figure 6.1: Ranges

[2]	14	16	3
[3]	17	17	1
[4]	19	23	5
[5]	27	31	5
[6]	28	32	5
[7]	29	33	5

**Inter-range** methods act on the collection of ranges as a whole. These include `disjoin`, `reduce`, `gaps`, and `range`. An illustration is `reduce`, which reduces overlapping ranges into a single range, as illustrated in the lower panel of Figure 6.1.

```
> reduce(ir)
```

```
IRanges of length 2
```

```
  start end width
```

```
[1]    7  18   12
```

```
[2]   22  28    7
```

`coverage` is an inter-range operation that calculates how many ranges overlap individual positions. Rather than returning ranges, `coverage` returns a compressed representation (run-length encoding)

```
> coverage(ir)
```

```
integer-Rle of length 28 with 12 runs
```

```
Lengths: 6 2 4 1 2 3 3 1 1 3 1 1
```

```
Values : 0 1 2 1 2 1 0 1 2 3 2 1
```

The run-length encoding can be interpreted as ‘a run of length 6 of nucleotides covered by 0 ranges, followed by a run of length 2 of nucleotides covered by 1 range...’.

**Between** methods act on two (or sometimes more) *IRanges* instances. These include `intersect`, `setdiff`, `union`, `pintersect`, `psetdiff`, and `punion`.

The `countOverlaps` and `findOverlaps` functions operate on two sets of ranges. `countOverlaps` takes its first argument (the `query`) and determines how many of the ranges in the second argument (the `subject`) each overlaps. The result is an integer vector with one element for each member of `query`. `findOverlaps` performs a similar operation but returns a more general matrix-like structure that identifies each pair of query / subject overlaps. Both arguments allow some flexibility in the definition of ‘overlap’.

Common operations on ranges are summarized in Table 6.2.



Table 6.2: Common operations on *IRanges*, *GRanges* and *GRangesList*.

Category	Function	Description
Accessors	<code>start, end, width</code>	Get or set the starts, ends and widths
	<code>names</code>	Get or set the names
	<code>mcols, metadata</code>	Get or set metadata on elements or object
	<code>length</code>	Number of ranges in the vector
Ordering	<code>range</code>	Range formed from min start and max end
	<code>&lt;, &lt;=, &gt;, &gt;=, ==, !=</code>	Compare ranges, ordering by start then width
	<code>sort, order, rank</code>	Sort by the ordering
	<code>duplicated</code>	Find ranges with multiple instances
Arithmetic	<code>unique</code>	Find unique instances, removing duplicates
	<code>r + x, r - x, r * x</code>	Shrink or expand ranges <code>r</code> by number <code>x</code>
	<code>shift</code>	Move the ranges by specified amount
	<code>resize</code>	Change width, anchoring on start, end or mid
Set operations	<code>distance</code>	Separation between ranges (closest endpoints)
	<code>restrict</code>	Clamp ranges to within some start and end
	<code>flank</code>	Generate adjacent regions on start or end
	<code>reduce</code>	Merge overlapping and adjacent ranges
Overlaps	<code>intersect, union, setdiff</code>	Set operations on reduced ranges
	<code>pintersect, punion, psetdiff</code>	Parallel set operations, on each <code>x[i]</code> , <code>y[i]</code>
	<code>gaps, pgap</code>	Find regions not covered by reduced ranges
	<code>disjoin</code>	Ranges formed from union of endpoints
Coverage	<code>findOverlaps</code>	Find all overlaps for each <code>x</code> in <code>y</code>
	<code>countOverlaps</code>	Count overlaps of each <code>x</code> range in <code>y</code>
	<code>nearest</code>	Find nearest neighbors (closest endpoints)
	<code>precede, follow</code>	Find nearest <code>y</code> that <code>x</code> precedes or follows
Extraction	<code>x %in% y</code>	Find ranges in <code>x</code> that overlap range in <code>y</code>
	<code>coverage</code>	Count ranges covering each position
	<code>r[i]</code>	Get or set by logical or numeric index
	<code>r[[i]]</code>	Get integer sequence from <code>start[i]</code> to <code>end[i]</code>
Split, combine	<code>subsetByOverlaps</code>	Subset <code>x</code> for those that overlap in <code>y</code>
	<code>head, tail, rev, rep</code>	Conventional R semantics
	<code>split</code>	Split ranges by a factor into a <i>RangesList</i>
	<code>c</code>	Concatenate two or more range objects

**mcols and metadata** The *GRanges* class (actually, most of the data structures defined or extending those in the *IRanges* package) has two additional very useful data components. The `mcols` function allows information on each range to be stored and manipulated (e.g., subset) along with the *GRanges* instance. The element metadata is represented as a *DataFrame*, defined in *IRanges* and acting like a standard *R data.frame* but with the ability to hold more complicated data structures as columns (and with element metadata of its own, providing an enhanced alternative to the *Biobase* class *AnnotatedDataFrame*).

```
> mcols(genes) <- DataFrame(EntrezId=c("42865", "2768869"),
+                             Symbol=c("kal-1", "CG34330"))
```

`metadata` allows addition of information to the entire object. The information is in the form of a list; any data can be provided.

```
> metadata(genes) <-
+   list(CreatedBy="A. User", Date=date())
```

***GRangesList*** The *GRanges* class is extremely useful for representing simple ranges. Some next-generation sequence data and genomic features are more hierarchically structured. A gene may be represented by several exons within it. An aligned read may be represented by discontinuous ranges of alignment to a reference. The *GRangesList* class represents this type of information. It is a list-like data structure, which each element of the list itself a *GRanges* instance. The gene FBgn0039155 contains several exons, and can be represented as a list of length 1, where the element of the list contains a *GRanges* object with 7 elements:

*GRangesList* of length 1:

\$FBgn0039155

*GRanges* with 7 ranges and 2 metadata columns:

	seqnames	ranges	strand	exon_id	exon_name
	<Rle>	<IRanges>	<Rle>	<integer>	<character>
[1]	chr3R	[19967117, 19967382]	+	45515	<NA>
[2]	chr3R	[19970915, 19971592]	+	45516	<NA>
[3]	chr3R	[19971652, 19971770]	+	45517	<NA>
[4]	chr3R	[19971831, 19972024]	+	45518	<NA>
[5]	chr3R	[19972088, 19972461]	+	45519	<NA>
[6]	chr3R	[19972523, 19972589]	+	45520	<NA>
[7]	chr3R	[19972918, 19973212]	+	45521	<NA>

---

seqlengths:

chr3R

27905053

The *GRangesList* object has methods one would expect for lists (e.g., `length`, sub-setting). Many of the methods introduced for working with *IRanges* are also available, with the method applied element-wise.

**The *GenomicFeatures* package** Many public resources provide annotations about genomic features. For instance, the UCSC genome browser maintains the ‘knownGene’ track of established exons, transcripts, and coding sequences of many model organisms. The *GenomicFeatures* package provides a way to retrieve, save, and query these resources. The underlying representation is as sqlite data bases, but the data are available in *R* as *GRangesList* objects. The following exercise explores the *GenomicFeatures* package and some of the functionality for the *IRanges* family introduced above.

## Exercise 9

Load the *TxDb.Dmelanogaster.UCSC.dm3.ensGene* annotation package, and create an alias `txdb` pointing to the *TranscriptDb* object this class defines.

Select just those elements corresponding to flybase gene ids *FBgn0002183*, *FBgn0003360*, *FBgn0025111*, and *FBgn0036449*. Use `reduce` to simplify gene models, so that exons that overlap are considered ‘the same’.

```
> library(TxDb.Dmelanogaster.UCSC.dm3.ensGene)
> txdb <- TxDb.Dmelanogaster.UCSC.dm3.ensGene # alias
> ex0 <- exonsBy(txdb, "gene")
> head(table(elementLengths(ex0)))

  1    2    3    4    5    6
3182 2608 2070 1628 1133  886

> ids <- c("FBgn0002183", "FBgn0003360", "FBgn0025111", "FBgn0036449")
> ex <- reduce(ex0[ids])
```

Most down-stream analysis of short read sequences is based on reads aligned to reference genomes.

```
> fl <- system.file("extdata", "ex1.sam", package="Rsamtools")
> strsplit(readLines(fl, 1), "\t")[[1]]
```

@DC?2.LFFGE@7>C62>BGABGB9HFBAFIIHEGFIIIHFAIIGDA<FC	Query base quals
AS:i:0	Alignment score
NM:i:0	Edit distance to reference
NH:i:10	Number of hits
XS:A:-	Strand
HI:i:0	Hit index for this alignment

50

Table 6.3: Fields in a SAM record. From <http://samtools.sourceforge.net/samtools.shtml>

Field	Name	Value
1	QNAME	Query (read) NAME
2	FLAG	Bitwise FLAG, e.g., strand of alignment
3	RNAME	Reference sequence NAME
4	POS	1-based leftmost POSition of sequence
5	MAPQ	MAPping Quality (Phred-scaled)
6	CIGAR	Extended CIGAR string
7	MRNM	Mate Reference sequence NaMe
8	MPOS	1-based Mate POSition
9	ISIZE	Inferred insert SIZE
10	SEQ	Query SEQUENCE on the reference strand
11	QUAL	Query QUALity
12+	OPT	OPTIONal fields, format TAG:VTYPE:VALUE

alignment is encoded in the ‘flag’ field. The alignment record also includes a measure of mapping quality, and a CIGAR string describing the nature of the alignment. In this case, the CIGAR is 36M, indicating that the alignment consisted of 36 Matches or mismatches, with no indels or gaps; indels are represented by I and D; gaps (e.g., from alignments spanning introns) by N.

BAM files encode the same information as SAM files, but in a format that is more efficiently parsed by software; BAM files are the primary way in which aligned reads are imported in to R.

**Aligned reads in R** The `readGappedAlignments` function from the *GenomicRanges* package reads essential information from a BAM file in to R. The result is an instance of the *GappedAlignments* class. The *GappedAlignments* class has been designed to allow useful manipulation of many reads (e.g., 20 million) under moderate memory requirements (e.g., 4 GB).

```
> alnFile <- system.file("extdata", "ex1.bam", package="Rsamtools")
> aln <- readGappedAlignments(alnFile)
> head(aln, 3)
```

GappedAlignments with 3 alignments and 0 metadata columns:

```
      seqnames strand      cigar  qwidth  start    end    width
      <Rle>   <Rle> <character> <integer> <integer> <integer> <integer>
[1]      seq1      +       36M        36        1       36       36
[2]      seq1      +       35M        35        3       37       35
[3]      seq1      +       35M        35        5       39       35
      ngap
      <integer>
[1]        0
[2]        0
[3]        0
---
seqlengths:
  seq1 seq2
1575 1584
```

## Parameters for scanning BAM files

### Description

Use `ScanBamParam()` to create a parameter object influencing what fields and which records are imported from a (binary) BAM file. Use of which requires that a BAM index file (<filename>.bai) exists.

The `readGappedAlignments` function takes an additional argument, `param`, allowing the user to specify regions of the BAM file (e.g., known gene coordinates) from which to extract alignments.

A *GappedAlignments* instance is like a data frame, but with accessors as suggested by the column names. It is easy to query, e.g., the distribution of reads aligning to each strand, the width of reads, or the cigar strings

```
> table(strand(aln))

+   -   *
1647 1624   0

> table(width(aln))

30  31  32  33  34  35  36  38  40
 2  21   1   8  37 2804 285   1 112

> head(sort(table(cigar(aln)), decreasing=TRUE))

      35M      36M      40M      34M      33M 14M4I17M
2804      283      112      37         6         4
```

### Exercise 10

Use *bigdata*, *file.path* and *dir* to obtain file paths to the BAM files. These are a subset of the aligned reads, overlapping just four genes.

Input the aligned reads from one file using *readGappedAlignments*. Explore the reads, e.g., using *table* or *xtabs*, to summarize which chromosome and strand the subset of reads is from.

The object *ex* created earlier contains coordinates of four genes. Use *countOverlaps* to first determine the number of genes an individual read aligns to, and then the number of uniquely aligning reads overlapping each gene. Since the RNAseq protocol was not strand-sensitive, set the strand of *aln* to *\**.

Write the sequence of steps required to calculate counts as a simple function, and calculate counts on each file. On Mac or Linux, can you easily parallelize this operation?

**Solution:** We discover the location of files using standard *R* commands:

```
> bigdata <- system.file("bigdata", package="SequenceAnalysisData")
> fls <- dir(file.path(bigdata, "bam"), ".bam$", full=TRUE) ##
> names(fls) <- sub("_.*", "", basename(fls))
```

Use *readGappedAlignments* to input data from one of the files, and standard *R* commands to explore the data.

```
> ## input
> aln <- readGappedAlignments(fls[1])
> xtabs(~seqnames + strand, as.data.frame(aln))

      strand
seqnames -   +
chr3L 5974 5402
chrX   2283 2278
```

To count overlaps in regions defined in a previous exercise, load the regions.

```
> data(ex)          # from an earlier exercise
```

Many RNA-seq protocols are not strand aware, i.e., reads align to the plus or minus strand regardless of the strand on which the corresponding gene is encoded. Adjust the strand of the aligned reads to indicate that the strand is not known.

```
> strand(aln) <- "*" # protocol not strand-aware
```

One important issue when counting reads is to make sure that the reference names in both the annotation and the read files are identical.

### Exercise 11

Check the reference name in both the `ex` and `aln`. If they are not similar, how could you correct them?

For simplicity, we are interested in reads that align to only a single gene. Count the number of genes a read aligns to...

```
> hits <- countOverlaps(aln, ex)
> table(hits)
```

```
hits
  0    1    2
772 15026 139
```

and reverse the operation to count the number of times each region of interest aligns to a uniquely overlapping alignment.

```
> cnt <- countOverlaps(ex, aln[hits==1])
```

A simple function for counting reads is

```
> counter <-
+   function(filePath, range)
+ {
+   aln <- readGappedAlignments(filePath)
+   strand(aln) <- "*"
+   hits <- countOverlaps(aln, range)
+   countOverlaps(range, aln[hits==1])
+ }
```

This can be applied to all files using `sapply`

```
> counts <- sapply(fls, counter, ex)
```

The counts in one BAM file are independent of counts in another BAM file. This encourages us to count reads in each BAM file in parallel, decreasing the length of time required to execute our program. On Linux and Mac OS, a straight-forward way to parallelize this operation is:

```
> if (require(parallel))
+   simplify2array(mclapply(fls, counter, ex))
```

**Flexible BAM parsing with [Rsamtools](#)** The *GappedAlignments* class inputs only some of the fields of a BAM file, and may not be appropriate for all uses. In these cases the `scanBam` function in [Rsamtools](#) provides greater flexibility. The idea is to view BAM files as a kind of data base. Particular regions of interest can be selected, and the information in the selection restricted to particular fields. These operations are determined by the values of a *ScanBamParam* object, passed as the named `param` argument to `scanBam`.

### Exercise 12

Consult the help page for *ScanBamParam*, and construct an object that restricts the information returned by a `scanBam` query to the aligned read DNA sequence. Your solution will use the `what` parameter to the *ScanBamParam* function.

Use the *ScanBamParam* object to query a BAM file, and calculate the GC content of all aligned reads. Summarize the GC content as a histogram (Figure 6.2).




Figure 6.2: GC content in aligned reads

**Solution:**

```
> gcFunction <-
+   function(x)
+ {
+   alf <- alphabetFrequency(x, as.prob=TRUE)
+   rowSums(alf[,c("G", "C")])
+ }
> param <- ScanBamParam(what="seq")
> seqs <- scanBam(fls[[1]], param=param)
> readGC <- gcFunction(seqs[[1]][["seq"]])
> hist(readGC)
```

**Part II**

**Differential Representation**



## Chapter 7

# RNA-seq Work Flows

### 7.1 Varieties of RNA-seq

RNA-seq experiments typically ask about differences in transcription of genes or other features across experimental groups. The analysis of designed experiments is statistical, and hence an ideal task for *R*. The overall structure of the analysis, with tens of thousands of features and tens of samples, is reminiscent of microarray analysis; some insights from the microarray domain will apply, at least conceptually, to the analysis of RNA-seq experiments.

The most straight-forward RNA-seq experiments quantify abundance for known gene models. The known models are derived from reference databases, reflecting the accumulated knowledge of the community responsible for the data. The ‘knownGenes’ track of the UCSC genome browser represents one source of such data. A track like this describes, for each gene, the transcripts and exons that are expected based on current data. The *GenomicFeatures* package allows ready access to this information by creating a local database out of the track information. This data base of known genes is coupled with high throughput sequence data by counting reads overlapping known genes and modeling the relationship between treatment groups and counts.

A more ambitious approach to RNA-seq attempts to identify novel transcripts. This requires that sequenced reads be assembled into contigs that, presumably, correspond to expressed transcripts that are then located in the genome. Regions identified in this way may correspond to known transcripts, to novel arrangements of known exons (e.g., through alternative splicing), or to completely novel constructs. We will not address the identification of completely novel transcripts here, but will instead focus on the analysis of the designed experiments: do the transcript abundances, novel or otherwise, differ between experimental groups?

### 7.2 Work flows and upstream analysis

RNA-seq work flows aim at measuring gene expression through assessment of mRNA abundance. Work flows involve:

1. Experimental design.
2. Wet-lab protocols for mRNA extraction and reverse transcription to cDNA.
3. Sequencing.
4. Alignment of sequenced reads to a reference genome.
5. Summarizing of the number of reads aligning to a region.
6. Normalization of samples to accommodate purely technical differences in preparation.
7. Statistical assessment of differential representation, including specification of an appropriate error model.
8. Interpretation of results in the context of original biological questions.

The inference is that higher levels of gene expression translate to more abundant cDNA, and greater numbers of reads aligned to the reference genome. The enumeration above seems simplistic, but oddly enough one has concerns and commentary on each point.

### 7.2.1 Experimental design

Obviously one should follow best practices for designing experiments appropriate for the data under analysis. A typical experiment will have one or several groups. Because there is uncertainty in each measurement, we require replication. Previous work shows that technical replication (repeating identical wet-lab and sequencing protocols on a single biological sample) introduces variation that is small [13] compared to biological replicates (using different samples). Most RNA-seq experiments require biological replication, and seldom include technical replicates.

How many biological replicates? It is helpful to think in terms of orders of magnitude – biological treatments with strong and consistent consequences for gene expression will be detected with a handful – 2 or 3 – replicates per treatment. Conversely, statistically subtle effects will not be much revealed by samples of say 5 or 8, but will instead require 10's or 100's of samples.

How complicated an experimental design? The advice must be to 'keep it simple'. There are many interesting biological questions that one could ask, but experimental designs with more than one or at most two factors, or with multiple levels per factor, will undermine statistical power and complicate analysis. There are exceptions of course, for instance a time course design or an experiment with two or more factors, but these require strong *a priori* motivation and confidence that the design is amenable to analysis even in the face of wet-lab or sequencing catastrophe.

What kind of treatment? Two 'lessons learned' from microarray analysis and applicable to RNA-seq inform this question. (a) It is necessary to normalize observations between samples to accommodate purely technical variation in overall patterns of expression. For example, samples provided to the sequencer have different amounts of DNA, resulting in variation in total numbers of sequenced and aligned reads independent of any difference in gene-level differential representation. This implies that the treatment should *affect only a fraction of the genes assayed*, otherwise treatment effects and protocol artifacts are confounded. (b) Between-gene measures of expression differ for reasons unrelated to levels of expression. For instance, standard protocols mean that a long gene is sequenced more often than a short gene, even when the number of mRNA molecules of the two genes are identical. This means that the most productive approach to differential representation will *compare genes across samples*, rather than compare levels of representation of different genes (gene set enrichment analysis and other approaches to between-gene comparison are statistically interesting in part because of the need to overcome between-gene differences arising for purely technical reasons). The combination of lessons (a) and (b) dictate that the treatment should affect only a subset of the genes under study, and that 'interesting' results correspond to treatment groups with differences at the gene level. *A priori* motivation, e.g., about well-defined pathways as targets of differential representation, may trump part (b) of this guideline.

The reality of executing designed experiments may mean that there are known but unavoidable factors that confound the analysis, but that are not of fundamental biological interest. Perhaps samples are being processed by different groups, or processing is spread over several months to accommodate personnel or sequencer availability. It is essential to avoid confounding such factors with biologically relevant parts of the experiment. Having acknowledged a potentially confounding factor, what is to be done? A first reaction might be randomization – arrange for samples to be processed in a random order, for instance, rather than by treatment group – but a better strategy is usually to include a blocking factor, e.g., processed by lab 'A' versus lab 'B' and to ensure that treatments are represented by replicates in each blocking factor. The down-stream analysis can then use replication to statistically accommodate such effects.

### 7.2.2 Wet-lab protocols, sequencing, and alignment

The important point here is that wet-lab protocols, sequencing reactions, and alignment introduce artifacts that need to be acknowledged and, if possible, accommodated in down-stream analysis. These artifacts and

Table 7.1: Statistical issues in RNA-seq differential expression.

Analysis stage	Issues
Summarizing	Counts versus RPKM and other summaries.
Normalization	Robust estimates of library size.
Differential expression	Appropriate error model (Negative Binomial, Poisson, . . . ); dispersion (under negative binomial) as parameter requiring estimation; ‘shrinkage’ to balance accuracy of per-gene estimates with precision of experiment-wide estimates.
Testing	Filtering to reduce multiple comparisons & false discovery rate.

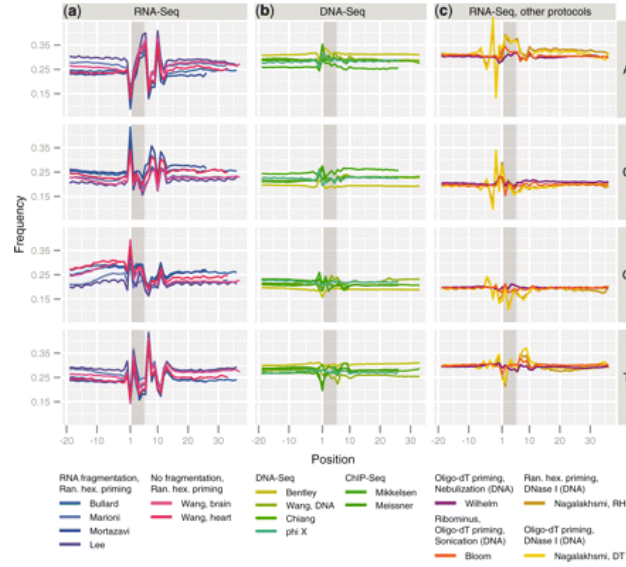


Figure 7.1: Nucleotide frequency versus position relative to start of alignment, various experiments and protocols; see [9].

approaches to their remediation are discussed in the following sections.

## 7.3 Statistical analysis

Important statistical issues are summarized in Table 7.1.

### 7.3.1 Summarizing

The Summarizing process tallies the number of reads aligning in each region (e.g., gene) of interest. The simplest method is to simply count reads overlapping each region, dividing by the length of the region of interest to accommodate differences in gene length. This is the ‘RPKM’ (reads per kilobase per million reads) of Mortazavi et al. [16]. One problem with this approach is that reads are not sampled uniformly across genes (Figure 7.1; [9]), so gene length (the ‘PK’ part of RPKM) is not a good proxy for expression level.

More fundamentally, each read represents an observation, and contributes to the certainty with which a gene is measured as ‘expressed’. A summary measure like RPKM fails to incorporate uncertainty – a particular value of RPKM might result from alignment of one or 100 reads. This contrasts with a simple count of the number of reads in the region of interest. Furthermore, count data has known statistical

properties that can be exploited in down-stream statistical analysis. Thus the result of summarization most useful for assessing differential expression is read count.

How to count? For instance, should a read that partly overlaps a 5' UTR or an intron be included in a tally? What about reads that overlap multiple genes? This is a non-trivial question because alignment is only approximate (reflecting sequencing and other biases) and because sample preparation protocols and organism biology (e.g., whether the UTR or fully mature mRNA is sequenced) may dictate particular counting strategies; more elaborate counting strategies might be entertained for paired end reads. Anders enumerates some counting strategies<sup>1</sup>; these are implemented in his *HTSeq* python scripts, in `summarizeOverlaps` in the *GenomicRanges* package, or in functions in the (linux-only) *Rsubread* and *gmapR* packages.

### 7.3.2 Normalization

Normalization arises from the need to correct for purely technical differences between samples. The most common symptom of the need for normalization is differences in the total number of aligned reads. The 'M' part of RPKM measure mentioned in the context of summarization is one way of normalizing for total count. This normalization is not appropriate, because the distribution of aligned reads across genes within a sample is not uniform – some regions receive many more alignments than do others – and this distribution may differ between samples.

The overall strategy with normalization is to choose an appropriate baseline, and express sample counts relative to that baseline. There are several approaches to choice of appropriate baseline. One might choose total count for normalization, but this is a poor choice when one or a few regions of interest are very well represented – we are normalizing to the well-represented genes rather than to sequencing depth in each sample. Other straight-forward approaches include use of house-keeping genes, or the expression level from a particular quantile of the distribution of gene expression values of each sample [4]. One might attempt a robust estimate of sample abundance that is less sensitive to extreme outliers, e.g., the trimmed geometric mean of counts [1]. Another approach is TMM [20], which measures the trimmed mean of M and A values (M values are the log fold change in the number of reads aligning to a region of interest measured relative to an average or arbitrary sample, A is the average count of a gene; the trimmed mean discards regions of interest that have extreme M or A values and calculates the mean M value of the remainder); the inverse of this mean is used to weight samples. More data-driven approaches exploiting the gene-specific properties include conditional quantile normalization (implemented in the *cqn* package; [10]).

Another approach to normalization, increasingly popular as experiment size and data consistency increases, is to perform a data transformation and apply normalization methods developed for analysis of microarrays. Examples of this approach include `varianceStabilizingTransformation` from the *DESeq* package, and `voom` from the *limma* package; see the corresponding help pages of these functions for details).

### 7.3.3 Error model

1. Negative binomial error model often appropriate – combination of Poisson ('shot' noise, technical variation in read counts) and variation between biological samples.
2. Requires estimate of 'dispersion'
3. 'Shrinkage' of gene-specific estimates to average.
4. *DESeq*, *edgeR* take different approaches, especially to dealing with outlier / extreme values.
5. Addition modeling approaches possible, e.g., implemented in the *DSS* package [26].

### 7.3.4 Multiple comparison

1. Increase statistical power and reduce false discovery rate by filtering regions of interest prior to analysis.

---

<sup>1</sup><http://www-huber.embl.de/users/anders/HTSeq/doc/count.html>

Table 7.2: Selected *Bioconductor* packages for RNA-seq analysis.

Package	Description
<a href="#"><i>EDASeq</i></a>	Exploratory analysis and QA; also <a href="#"><i>qrqc</i></a> , <a href="#"><i>ShortRead</i></a> , <a href="#"><i>DESeq</i></a> .
<a href="#"><i>edgeR</i></a> , <a href="#"><i>DESeq</i></a>	Generalized Linear Models using negative binomial error.
<a href="#"><i>BitSeq</i></a>	Bayesian inference of individual transcript abundances followed by differential expression.
<a href="#"><i>DEXSeq</i></a>	Exon-level differential representation.
<a href="#"><i>DSS</i></a> , <a href="#"><i>vsn</i></a> , <a href="#"><i>cqn</i></a>	RNA-seq normalization methodologies. Also <i>voom</i> in <a href="#"><i>limma</i></a> .
<a href="#"><i>goseq</i></a>	Gene set enrichment tailored to RNAseq count data; also <a href="#"><i>limma</i></a> 's <i>roast</i> or <i>camera</i> after transformation with <i>voom</i> .
<a href="#"><i>easyRNASeq</i></a>	Workflow; also <a href="#"><i>ArrayExpressHTS</i></a> , <a href="#"><i>rnaSeqMap</i></a> , <a href="#"><i>oneChannelGUI</i></a> .
<a href="#"><i>Rsubread</i></a>	Alignment (Linux only); also <a href="#"><i>gmapR</i></a> ; <a href="#"><i>Biostrings</i></a> <i>matchPDict</i> for special-purpose alignments.
<a href="#"><i>cummeRbund</i></a>	Exploration and analysis of Cufflinks results.

- Motivation (a): just because genes are assayed does not mean, *a priori*, that they represent something requiring a statistical test. (b) Some observations, e.g., zero counts across all samples, cannot possibly be statistically significant, independent of hypothesis under investigation.
- Approach – detection or ‘*K* over *A*’-style filter; representation of a minimum of *A* (normalized) read counts in at least *K* samples. *A* usually measured as counts per million. Guidelines for choice of values a little *ad hoc*; see, e.g., the [\*edgeR\*](#) user manual. Variance filter, e.g., IQR (inter-quartile range) provides a robust estimate of variability; can be used to rank and discard least-varying regions.

### 7.3.5 *Bioconductor* software

*Bioconductor* packages play a role in several stages of an RNA-seq analysis (Table 7.2; a more comprehensive list is under the [RNAseq](#) and [HighThroughputSequencing](#) BiocViews terms). The [\*GenomicRanges\*](#) infrastructure can be effectively employed to quantify known exon or transcript abundances. Quantified abundances are in essence a matrix of counts, with rows representing features and columns samples. The [\*edgeR\*](#) [20] and [\*DESeq\*](#) [1] packages facilitate analysis of this data in the context of designed experiments, and are appropriate when the questions of interest involve between-sample comparisons of relative abundance. The [\*DEXSeq\*](#) package extends the approach in [\*edgeR\*](#) and [\*DESeq\*](#) to ask about within-gene, between group differences in exon use, i.e., for a given gene, do groups differ in their exon use?

## Chapter 8

# *DESeq* Work Flow Exercises

This chapter walks through an RNA-seq differential expression work flow. It is based on the *DESeq* vignette, reproduced in the Appendix A. Anders and Huber [1] provide additional information.

### Exercise 13

- Visit the *DESeq* landing page<sup>1</sup>.
- Retrieve a copy of the ‘Analyzing RNA-Seq data with the “*DESeq*” package’ vignette pdf.
- Retrieve a copy of the R script corresponding to the vignette.
- How can the `vignette` function be used to access the vignette and script from within R, without needing to visit the Bioconductor web site?

## 8.1 Data input and preparation

### Exercise 14

- Read through sections 1 and 2 of the *DESeq* vignette.
- Evaluate the code chunks (numbered 1 through 16) corresponding to sections 1 and 2. Step through these one at a time, reflecting on the commands that you are evaluating and the decisions that are being made about your data.
- Summarize key steps that have been taken to this point, emphasizing analysis decisions you have made, the assumptions underlying these decisions, and the possible consequences of an incorrect decision.

## 8.2 Inference

### Exercise 15

- Read through section 3.1 of the *DESeq* vignette.
- Evaluate the code chunks (17 through 28) corresponding to section 3.1.
- Summarize key steps that have been taken to this point, emphasizing analysis decisions you have made, the assumptions underlying these decisions, and the possible consequences of an incorrect decision.

---

<sup>1</sup><http://bioconductor.org/packages/release/bioc/html/DESeq.html>

## 8.3 Independent filtering

### Exercise 16

- a. Read through section 5 of the [DESeq](#) vignette.
- b. Evaluate the code chunks corresponding to section 5.
- c. How do you know when independent filtering is appropriate and effective?

## 8.4 Data quality assessment

### 8.4.1 Preliminary transformation

Many of the data quality challenges associated with RNA-seq differential expression data are similar to those seen in microarray studies. A convenient approach to data assessment is then to transform the count data to an approximately appropriate scale, and to use transformed data and microarray quality assessment best practices to guide RNA-seq differential expression best practices. The [DESeq](#) vignette adopts this approach, using the variance stabilizing transformation outlined in section 6.

### Exercise 17

- a. Read through the introductory part of section 6 of the [DESeq](#) vignette to gain at least a superficial understanding of the variance stabilizing transformation.

### 8.4.2 Quality assessment

### Exercise 18

- a. Read through section 7 of the [DESeq](#) vignette.
- b. Evaluate the code chunks corresponding to sections 7.1 to 7.3.
- c. Compare the two panels of the vignette Figure 15. What is the importance of data transformation prior to quality assessment?

## 8.5 Frequently asked questions

**No or incomplete replicates** See sections 3.2 and 3.3 of the [DESeq](#) vignette.

**Complicated designs** We discuss complicated designs in Section ??.

**Different results from different analyses** Making different decisions can lead to different results. Which result is correct? The underlying decision might involve a single package (e.g., choice of dispersion estimates) or might arise from choice of packages used in an analysis (e.g., use of [DESeq](#) versus [edgeR](#)). Unfortunately, there is no straight-forward answer to which analysis is ‘correct’. The appropriate analysis is the one that makes assumptions that most closely match the nature of the individual data set, or for which results are relatively robust to violations of the assumptions of the analysis.

**Any rules of thumb to help identify a robust analysis?**

# Part III

## Variant Calls



# Chapter 9

## Variant Work Flows

### 9.1 Variants

#### 9.1.1 Varieties of variant-related work flows

- Single nucleotide polymorphism
- Copy number change
- Structural variation
- Long-range interaction

#### 9.1.2 Work flows

- Alignment. requires tools sensitive to variation, e.g., *GSNAP*, *BWA*; *Bowtie* not optimized for this.
- Variant calling, e.g., *GATK*<sup>1</sup>.
- Filter.
- Biological context – variant annotation.
- Integrative analysis, e.g., GWAS, genetical genomics.

#### 9.1.3 *Bioconductor* software

Selected *Bioconductor* software relevant to DNA-seq work flows are summarized in Table 9.1. The *gmapR* package provides access to the well-respected *GSNAP* aligner. *VariantTools* is an emerging tool that works with aligned reads to call single-sample and sample-specific variants; we will work with *VariantTools* as part of this course. Packages such as *cn.mops* and *exomeCopy* identify copy number variants from high-throughput sequence data; *r3Cseq* provides tools to analyze and visualize long-range genomic interactions.

*VariantAnnotation* provides facilities for manipulated called variants stored in VCF files; the facilities are very flexible, including simple range-based filtering, look-up in dbSNP, and coding and effect prediction using standard data bases. *VariantAnnotation* plays well with *ensemblVEP* to feed data through the Ensembl Variant Effect Predictor perl script, and to *Bioconductor* facilities for SNP analysis and genetical genomics like *snpStats* and *GGtools*.

---

<sup>1</sup><http://www.broadinstitute.org/gatk/>

Table 9.1: Selected *Bioconductor* packages for DNA-seq analysis.

Package	Description
<i>gmapR</i>	Alignment (Linux only)
<i>VariantTools</i>	Single-sample and tumor specific variant calls
<i>deepSNV</i>	Sub-clonal SNVs in deep sequencing experiments
<i>cn.mops</i>	Mixture of Poissons copy number variation estimates
<i>exomeCopy</i>	Hidden Markov copy number variation estimates
<i>r3Cseq</i>	Long-range genomic interactions
<i>VariantAnnotation</i>	Manipulating and annotating VCF files
<i>ensemblVEP</i>	Interface to the Ensembl Variant Effect Predictor
<i>snpStats</i>	Down-stream GWAS; also <i>GWAStools</i> , <i>GGtools</i>

## 9.2 *VariantTools*

**WARNING:** This portion of the course is only available on Linux. The *VariantTools* package represents a work in progress.

### 9.2.1 Example data: lung cancer cell lines

#### Exercise 19

Load the *LungCancerLines* experiment data package. Check out the package description (e.g., `help(package="LungCancerLines")` for a brief description of the data. Use tools from previous exercises to explore the FASTQ and BAM files that are available in this package.

Open the *VariantTools* package vignette and associated script.

**Solution:** Here we load the package and discover paths to the BAM and FASTQ files, exploration is up to you!

```
> library(LungCancerLines)
> library(ShortRead)
> FastqFileList(LungCancerFastqFiles())
```

```
FastqFileList of length 4
names(4): H1993.first H1993.last H2073.first H2073.last
```

```
> LungCancerBamFiles()
```

```
BamFileList of length 2
names(2): H1993 H2073
```

Open the *VariantTools* vignette with

```
> vignette(package="VariantTools", "VariantTools")
```

or by visiting the package landing page<sup>2</sup>.

---

<sup>2</sup><http://bioconductor.org/packages/devel/bioc/html/VariantTools.html>

## 9.2.2 Calling single-sample variants

### Exercise 20

Create the data objects `p53`, `bams`, `bam`, and `tally.params` in section 2.1 of the vignette. In particular, consult the help page `?VariantTallyParam` to gain insight on the parameters that you have established.

Follow section 2.2 of the [VariantTools](#) vignette, paying particular attention to each function and the return value.

### Exercise 21

Filter variants based on standard quality filters, as in section 2.3 of the vignette.

Are there additional filters that might be appropriate at this stage of the analysis? Implement and additional filter following the hints in the brief section 2.4.

### Exercise 22

Return to the bottom of page 2 of the [VariantTools](#) package, and identify variants overlapping exons in the `p53` gene.

### Exercise 23

As a simple exercise, export called variants to a VCF file following section 4 of the vignette. Spend the time to discover the relationship between the variants you have identified, their representation in R, and their representation on disk.

## 9.2.3 Additional work flows

### Exercise 24

As time permits, explore sections of the vignette addressing variants called across samples

The previous exercise called variants in the metastatic sample, `H1993`. Conduct a similar analysis for the tumor sample `H2073`.

Explore section 3.1, calling sample-specific variants.

## Chapter 10

# Working with Called Variants

A major product of DNaseq experiments are catalogs of called variants (e.g., SNPs, indels). We will use the *VariantAnnotation* package to explore this type of data. Sample data included in the package are a subset of chromosome 22 from the *1000 Genomes* project. Variant Call Format (VCF; [full description](#)) text files contain meta-information lines, a header line with column names, data lines with information about a position in the genome, and optional genotype information on samples for each position.

Important operations on VCF files available with the *VariantAnnotation* package are summarized in Table 10.1.

### 10.1 Variant call format (VCF) files with *VariantAnnotation*

#### 10.1.1 Data input

##### Exercise 25

*The objective of this exercise is to compare the quality of called SNPs that are located in dbSNP, versus those that are novel.*

*Locate the sample data in the file system. Explore the metadata (information about the content of the file) using `scanVcfHeader`. Discover the ‘info’ fields VT (variant type), and RSQ (genotype imputation quality).*

*Input the sample data using `readVcf`. You’ll need to specify the genome build (`genome="hg19"`) on which the variants are annotated. Take a peak at the `rowData` to see the genomic locations of each variant.*

*Data resource often adopt different naming conventions for sequences. For instance, dbSNP uses abbreviations such as `ch22` to represent chromosome 22, whereas our VCF file uses `22`. Use `rowData` and `seqlevels<-` to extract the row data of the variants, and rename the chromosomes.*

**Solution:** Explore the header:

```
> library(VariantAnnotation)
> fl <- system.file("extdata", "chr22.vcf.gz", package="VariantAnnotation")
> (hdr <- scanVcfHeader(fl))

class: VCFHeader
samples(5): HG00096 HG00097 HG00099 HG00100 HG00101
meta(1): fileformat
fixed(1): ALT
info(22): LDAF AVGPOST ... VT SNPSOURCE
geno(3): GT DS GL

> info(hdr)[c("VT", "RSQ"),]
```

Table 10.1: Working with *VCF* files and data.

Category	Function	Description
Read	<code>scanVcfHeader</code>	Retrieve file header information
	<code>scanVcfParam</code>	Select fields to read in
	<code>readVcf</code>	Read VCF file into a VCF class
	<code>scanVcf</code>	Read VCF file into a list
Filter	<code>filterVcf</code>	Filter a VCF from one file to another
Write	<code>writeVcf</code>	Write a VCF file to disk
Annotate	<code>locateVariants</code>	Identify where variant overlaps a gene annotation
	<code>predictCoding</code>	Amino acid changes for variants in coding regions
	<code>summarizeVariants</code>	Summarize variant counts by sample
SNPs	<code>genotypeToSnpMatrix</code>	Convert genotypes to a SnpMatrix
	<code>GLtoGP</code>	Convert genotype likelihoods to genotypes
	<code>snpSummary</code>	Counts and distribution statistics for SNPs
Manipulate	<code>expand</code>	Convert CompressedVCF to ExpandedVCF
	<code>cbind, rbind</code>	Combine by column or row

DataFrame with 2 rows and 3 columns

	Number	Type	Description
	<character>	<character>	<character>
VT	1	String	indicates what type of variant the line represents
RSQ	1	Float	Genotype imputation quality from MaCH/Thunder

Input the data and peak at their locations:

```
> (vcf <- readVcf(fl, "hg19"))
```

```
class: CollapsedVCF
```

```
dim: 10376 5
```

```
rowData(vcf):
```

```
GRanges with 5 metadata columns: paramRangeID, REF, ALT, QUAL, FILTER
```

```
info(vcf):
```

```
DataFrame with 22 columns: LDAF, AVGPOST, RSQ, ERATE, THETA, CIEND, CIPOS, END, HOMLEN, HOMSEQ, SVLEN,
```

```
info(header(vcf)):
```

	Number	Type	Description
LDAF	1	Float	MLE Allele Frequency Accounting for LD
AVGPOST	1	Float	Average posterior probability from MaCH/Thunder
RSQ	1	Float	Genotype imputation quality from MaCH/Thunder
ERATE	1	Float	Per-marker Mutation rate from MaCH/Thunder
THETA	1	Float	Per-marker Transition rate from MaCH/Thunder
CIEND	2	Integer	Confidence interval around END for imprecise va...
CIPOS	2	Integer	Confidence interval around POS for imprecise va...
END	1	Integer	End position of the variant described in this r...
HOMLEN	.	Integer	Length of base pair identical micro-homology at...
HOMSEQ	.	String	Sequence of base pair identical micro-homology ...
SVLEN	1	Integer	Difference in length between REF and ALT alleles
SVTYPE	1	String	Type of structural variant
AC	.	Integer	Alternate Allele Count
AN	1	Integer	Total Allele Count
AA	1	String	Ancestral Allele, ftp://ftp.1000genomes.ebi.ac....
AF	1	Float	Global Allele Frequency based on AC/AN

```
[ reached getOption("max.print") -- omitted 6 rows ]
geno(vcf):
  SimpleList of length 3: GT, DS, GL
geno(header(vcf)):
  Number Type   Description
  GT 1      String Genotype
  DS 1      Float  Genotype dosage from MaCH/Thunder
  GL .      Float  Genotype Likelihoods

> head(rowData(vcf), 3)

GRanges with 3 ranges and 5 metadata columns:
      seqnames          ranges strand | paramRangeID
      <Rle>          <IRanges> <Rle> |      <factor>
rs7410291      22 [50300078, 50300078] * |      <NA>
rs147922003    22 [50300086, 50300086] * |      <NA>
rs114143073    22 [50300101, 50300101] * |      <NA>
      REF          ALT      QUAL      FILTER
      <DNAStringSet> <DNAStringSetList> <numeric> <character>
rs7410291          A          G      100      PASS
rs147922003        C          T      100      PASS
rs114143073        G          A      100      PASS
---
seqlengths:
22
NA
```

Rename chromosome levels:

```
> seqlevels(vcf, force=TRUE) <- c("22"="ch22")
```

## 10.2 SNP Annotation

Variants can be easily identified according to region such as `coding`, `intron`, `intergenic`, `spliceSite` etc. Amino acid coding changes are computed for the non-synonymous variants. SIFT and PolyPhen databases provide predictions of how severely the coding changes affect protein function. Additional annotations are easily crafted using the [GenomicRanges](#) and [GenomicFeatures](#) software in conjunction with *Bioconductor* and broader community annotation resources.

### Exercise 26

The [SNPlocs.Hsapiens.dbSNP.20101109](#) contains information about SNPs in a particular build of dbSNP. Load the package.

Review the following short helper function to query whether SNPs are in the data base (a version of this function will be introduced to [VariantAnnotation](#) before the next release)

```
> .isInDbSNP <-
+   function(vcf, seqname, rsid=TRUE)
+   {
+     snpLocs <- getSNPlocs(seqname)
+     idx <-      # correct seqname, width of variant == 1
+       ((seqnames(vcf) == seqname) & (width(rowData(vcf)) == 1L))
```

Figure 10.1: Quality scores of variants in dbSNP, compared to those not in dbSNP.

Table 10.2: Variant locations

Location	Details
coding	Within a coding region
fiveUTR	Within a 5' untranslated region
threeUTR	Within a 3' untranslated region
intron	Within an intron region
intergenic	Not within a transcript associated with a gene
spliceSite	Overlaps any of the first or last 2 nucleotides of an intron

```
+ idx <- as.vector(idx)
+ snps <- rowData(vcf)[idx]
+ result <- rep(NA, nrow(vcf))
+ result[idx] <- if (rsid) {
+   sub("rs", "", names(snps)) %in% snpLocs[["RefSNP_id"]]
+ } else {
+   start(snps) %in% snpLocs[["loc"]]
+ }
+ result
+ }
```

Create a data frame containing the dbSNP membership status and imputation quality of each SNP. Create a density plot to illustrate the results.

**Solution:** Discover whether SNPs are located in dbSNP, using our helper function.

```
> library(SNPlocs.Hsapiens.dbSNP.20101109)
> inDbSNP <- .isInDbSNP(vcf, "ch22")
> table(inDbSNP)
```

Create a data frame summarizing SNP quality and dbSNP membership:

```
> metrics <-
+   data.frame(inDbSNP=inDbSNP, RSQ=info(vcf)$RSQ)
```

Finally, visualize the data, e.g., using *ggplot2* (Figure 10.1).

```
> library(ggplot2)
> ggplot(metrics, aes(RSQ, fill=inDbSNP)) +
+   geom_density(alpha=0.5) +
+   scale_x_continuous(name="MaCH / Thunder Imputation Quality") +
+   scale_y_continuous(name="Density") +
+   theme(legend.position="top")
```

**Locating variants in and around genes** Variant location with respect to genes can be identified with the `locateVariants` function. Regions are specified in the `region` argument and can be one of the following constructors: `CodingVariants()`, `IntronVariants()`, `FiveUTRVariants()`, `ThreeUTRVariants()`, `IntergenicVariants()`, `SpliceSiteVariants()`, or `AllVariants()`. Location definitions are shown in Table 10.2.

## Exercise 27

Load the `TxDb.Hsapiens.UCSC.hg19.knownGene` annotation package, and read in the `chr22.vcf.gz` example file from the `VariantAnnotation` package.

Remembering to re-name sequence levels, use the `locateVariants` function to identify coding variants.

Summarize aspects of your data, e.g., did any coding variants match more than one gene? How many coding variants are there per gene ID?

**Solution:** Here we open the known genes data base, and read in the VCF file.

```
> library(TxDb.Hsapiens.UCSC.hg19.knownGene)
> txdb <- TxDb.Hsapiens.UCSC.hg19.knownGene
> fl <- system.file("extdata", "chr22.vcf.gz", package="VariantAnnotation")
> vcf <- readVcf(fl, "hg19")
> seqlevels(vcf, force=TRUE) <- c("22"="chr22")
```

The next lines locate coding variants.

```
> rd <- rowData(vcf)
> loc <- locateVariants(rd, txdb, CodingVariants())
> head(loc, 3)
```

GRanges with 3 ranges and 7 metadata columns:

	seqnames	ranges	strand	LOCATION	QUERYID	TXID
	<Rle>	<IRanges>	<Rle>	<factor>	<integer>	<integer>
[1]	chr22	[50301422, 50301422]	*	coding	24	73482
[2]	chr22	[50301476, 50301476]	*	coding	25	73482
[3]	chr22	[50301488, 50301488]	*	coding	26	73482

	CDSID	GENEID	PRECEDEID	FOLLOWID
	<integer>	<character>	<character>	<character>
[1]	217009	79087	<NA>	<NA>
[2]	217009	79087	<NA>	<NA>
[3]	217009	79087	<NA>	<NA>

---

seqlengths:

```
chr22
NA
```

To answer gene-centric questions data can be summarized by gene regardless of transcript.

```
> ## Did any coding variants match more than one gene?
> splt <- split(loc$GENEID, loc$QUERYID)
> table(sapply(splt, function(x) length(unique(x)) > 1))
```

```
FALSE  TRUE
  956    15
```

```
> ## Summarize the number of coding variants by gene ID
> splt <- split(loc$QUERYID, loc$GENEID)
> head(sapply(splt, function(x) length(unique(x))), 3)
```

```
113730  1890  23209
      22    15     30
```



**Amino acid coding changes** `predictCoding` computes amino acid coding changes for non-synonymous variants. Only ranges in `query` that overlap with a coding region in `subject` are considered. Reference sequences are retrieved from either a `BSgenome` or fasta file specified in `seqSource`. Variant sequences are constructed by substituting, inserting or deleting values in the `varAllele` column into the reference sequence. Amino acid codes are computed for the variant codon sequence when the length is a multiple of 3.

The `query` argument to `predictCoding` can be a `GRanges` or `VCF`. When a `GRanges` is supplied the `varAllele` argument must be specified. In the case of a `VCF` object, the alternate alleles are taken from `alt(<VCF>)` and the `varAllele` argument is not specified.

The result is a modified `query` containing only variants that fall within coding regions. Each row represents a variant-transcript match so more than one row per original variant is possible.

```
> library(BSgenome.Hsapiens.UCSC.hg19)
> coding <- predictCoding(vcf, txdb, seqSource=Hsapiens)
> coding[5:9]
```

GRanges with 5 ranges and 17 metadata columns:

	seqnames	ranges	strand	paramRangeID	
	<Rle>	<IRanges>	<Rle>	<factor>	
22:50301584	chr22	[50301584, 50301584]	-	<NA>	
	REF	ALT	QUAL	FILTER	
	<DNAStringSet>	<DNAStringSetList>	<numeric>	<character>	
22:50301584	C	T	100	PASS	
	varAllele	CDSLOC	PROTEINLOC	QUERYID	
	<DNAStringSet>	<IRanges>	<CompressedIntegerList>	<integer>	
22:50301584	A	[777, 777]	259	28	
	TXID	CDSID	GENEID	CONSEQUENCE	REFCODON
	<character>	<integer>	<character>	<factor>	<DNAStringSet>
22:50301584	73482	217009	79087	synonymous	CCG
	VARCODON	REFAA	VARAA		
	<DNAStringSet>	<AAStringSet>	<AAStringSet>		
22:50301584	CCA	P	P		

[ reached getOption("max.print") -- omitted 4 rows ]

---

seqlengths:

chr22

NA

Using variant rs114264124 as an example, we see `varAllele` A has been substituted into the `refCodon` CCG to produce `varCodon` CAG. The `refCodon` is the sequence of codons necessary to make the variant allele substitution and therefore often includes more nucleotides than indicated in the range (i.e. the range is 50302962, 50302962, width of 1). Notice it is the second position in the `refCodon` that has been substituted. This position in the codon, the position of substitution, corresponds to genomic position 50302962. This genomic position maps to position 698 in coding region-based coordinates and to triplet 233 in the protein. This is a non-synonymous coding variant where the amino acid has changed from R (Arg) to Q (Gln).

When the resulting `varCodon` is not a multiple of 3 it cannot be translated. The consequence is considered a `frameshift` and `varAA` will be missing.

```
> coding[coding$CONSEQUENCE == "frameshift"]
```

GRanges with 1 range and 17 metadata columns:

	seqnames	ranges	strand	paramRangeID
	<Rle>	<IRanges>	<Rle>	<factor>
22:50317001	chr22	[50317001, 50317001]	+	<NA>

```

          REF          ALT          QUAL          FILTER
    <DNASet> <DNASetList> <numeric> <character>
22:50317001      G      GCACT      233      PASS
      varAllele      CDSLOC      PROTEINLOC      QUERYID
    <DNASet> <IRanges> <CompressedIntegerList> <integer>
22:50317001      GCACT [808, 808]      270      359
      TXID      CDSID      GENEID CONSEQUENCE      REFCODON
    <character> <integer> <character>      <factor> <DNASet>
22:50317001      72592      214765      79174 frameshift      GCC
      VARCHODON      REFAA      VARAA
    <DNASet> <AASet> <AASet>
22:50317001      ACC      A
---
seqlengths:
chr22
NA

```

**SIFT and PolyPhen databases** From `predictCoding` we identified the amino acid coding changes for the non-synonymous variants. For this subset we can retrieve predictions of how damaging these coding changes may be. SIFT (Sorting Intolerant From Tolerant) and PolyPhen (Polymorphism Phenotyping) are methods that predict the impact of amino acid substitution on a human protein. The SIFT method uses sequence homology and the physical properties of amino acids to make predictions about protein function. PolyPhen uses sequence-based features and structural information characterizing the substitution to make predictions about the structure and function of the protein.

Collated predictions for specific dbSNP builds are available as downloads from the SIFT and PolyPhen web sites. These results have been packaged into *SIFT.Hsapiens.dbSNP132.db* and *PolyPhen.Hsapiens.dbSNP131.db* and are designed to be searched by rsid. Variants that are in dbSNP can be searched with these database packages. When working with novel variants, SIFT and PolyPhen must be called directly. See references for home pages.

Identify the non-synonymous variants and obtain the rsids.

```

> nms <- names(coding)
> idx <- coding$CONSEQUENCE == "nonsynonymous"
> nonsyn <- coding[idx]
> names(nonsyn) <- nms[idx]
> rsids <- unique(names(nonsyn)[grep("rs", names(nonsyn), fixed=TRUE)])

```

Detailed descriptions of the database columns can be found with `?SIFTdbColumns` and `?PolyPhenDbColumns`. Variants in these databases often contain more than one row per variant. The variant may have been reported by multiple sources and therefore the source will differ as well as some of the other variables.

```

> library(SIFT.Hsapiens.dbSNP132)
> ## rsids in the package
> head(keys(SIFT.Hsapiens.dbSNP132), 3)

[1] "rs10000692" "rs10001580" "rs10002700"

> ## list available columns
> cols(SIFT.Hsapiens.dbSNP132)

[1] "RSID"          "PROTEINID"     "AACHANGE"      "METHOD"        "AA"
[6] "PREDICTION"    "SCORE"         "MEDIAN"        "POSTIONSEQS"   "TOTALSEQS"

```

```

> ## select a subset of columns
> ## a warning is thrown when a key is not found in the database
> subst <- c("RSID", "PREDICTION", "SCORE", "AACHANGE", "PROTEINID")
> sift <- select(SIFT.Hsapiens.dbSNP132, keys=rsids, cols=subst)
> head(sift, 3)

```

```

      RSID PROTEINID AACHANGE PREDICTION SCORE
1 rs114264124 NP_077010   R233Q  TOLERATED  0.59
2 rs114264124 NP_077010   R233Q  TOLERATED  1.00
3 rs114264124 NP_077010   R233Q  TOLERATED  0.20

```

PolyPhen provides predictions using two different training datasets and has considerable information about 3D protein structure. See `?PolyPhenDbColumns` or the PolyPhen web site listed in the references for more details.

## 10.3 Large-scale filtering

One source of VCF files is from whole-genome sequencing and variant calling; an example data set is from Complete Genomics, and a subset is available on the Amazon Machine instance accompanying this course

```

> vcfDir <- "~/Seattle-Feb-2013"
> dir(vcfDir)
> vcfFile <- dir(vcfDir, ".gz$", full=TRUE)

```

The data were retrieved from the Complete Genomics web site, the first 750,000 of about 14 million variants selected, and the resulting file compressed and indexed. Indexing makes them accessible to fast queries using `which` argument to `ScanVcfParam`.

### Exercise 28

*The objective of this exercise is to filter the larger VCF file to a subset of interesting variants that we might wish to study in depth at a later date. We use the `filterVcf` function of the [VariantAnnotation](#) package to perform the filtering.*

*Start by taking a look at the (complicated) header information*

```

> hdr <- scanVcfHeader(vcfFile)

```

We'll be paying attention to the `SS INFO` field, and the `AD GENO` field. Determine the data types and possible values for these fields, using commands like

```

> info(hdr)
> geno(hdr)

```

What data type would you use to represent the `SS` field for a single or several VCF records in R? What about the `AD` field, across all samples?

This VCF file is big. While we can read this into memory all at once, we will often want to 'chunk' through a file, reading many (e.g., a million) records at a time. We do this by creating a `TabixFile` and specifying a `yieldSize` representing the size of the file that we'd like to read at each go. Create a `TabixFile` with `yieldSize=100000` and verify with a simple loop that the entire file appears to be processed in chunks of the specified size, along the lines of...

```

> tbx <- TabixFile(vcfFile, yieldSize=100000)
> open(tbx)
> while (len <- nrow(readVcf(tbx, "hg19")))
+   cat("read", len, "rows\n")

```

As you can see from the chunking exercise, it takes quite a bit of time to process these lines. To filter 14 million variants effectively, it can pay to do a cheaper ‘pre-filter’. Specifically, we’re interested in variants tagged ‘Germline’. If we were to represent each VCF record as an element of a character vector `x`, then we could write a one-liner that returned `TRUE` if the line contained the word ‘Germline’:

```
> grepl("Germline", x, fixed=TRUE)
```

This would be fast to read in to R, and fast to perform the filter. The `filterVcf` function allows us to specify a pre-filter that works just like this. The filters are constructed using the `FilterRules` constructor in [IRanges](#), by translating our one-liner into a simple function call, and placing the function call into a list.

```
> isGermline <- function(x)
+   grepl("Germline", x, fixed=TRUE)
> filters <- FilterRules(list(isGermline=isGermline))
```

The idea is that several filters are chained together. Each filter returns a logical vector indicating the subset of data to be processed by the next filter.

Here is our pre-filter in action:

```
> destination <- tempfile()           # temporary location
> filterVcf(vcfFile, "hg19", destination, prefilters=filters)
```

This is pretty fast, and drops the number of variants under consideration quite substantially, to about 110000.

Our next filter is more challenging to write. We’re interested in allelic depth, a summary of the evidence for the variant summarized in the `AD` `GENO` field. There are many variants, each sample has two values of `AD`, and there are two samples. This means that `AD` is a three-dimensional array. Our filter criteria is that the ratio of (‘alternate allele’ of the tumor sample or the ‘reference allele’ of the the normal sample) to total reads is greater than 0.1. Here’s function implementing this:

```
> allelicDepth <- function(x)
+ {
+   ## ratio of AD of the 'alternate allele' for the tumor sample
+   ## OR 'reference allele' for normal samples to total reads for
+   ## the sample should be greater than some threshold (say 0.1,
+   ## that is: at least 10% of the sample should have the allele
+   ## of interest)
+   ad <- geno(x)[["AD"]]
+   tumorPct <- ad[,1,2,drop=FALSE] / rowSums(ad[,1,,drop=FALSE])
+   normPct <- ad[,2,1, drop=FALSE] / rowSums(ad[,2,,drop=FALSE])
+   test <- (tumorPct > 0.1) | (normPct > 0.1)
+   !is.na(test) & test
+ }
```

We can add it to our list of filters

```
> filters <- FilterRules(list(isGermline=isGermline,
+                             allelicDepth=allelicDepth))
```

To use this filter, we actually need to fully parse the VCF file into the VCF instance; the pre-filter trick used for germline filtering is not enough. `filterVcf` allows us to perform a filter on the VCF instance, too, and does so after pre-filtering

```
> destination <- tempfile()
> filterVcf(vcfFile, "hg19", destination, prefilters=filters[1],
+           filters=filters[2])
```

*And finally input our interesting variants, confirming that we've done our filtering as desired.*

```
> vcf <- readVcf(destination, "hg19")  
> all(info(vcf)$SS == "Germline")  
> table(allelicDepth(vcf))
```

## Part IV

# Annotation and Visualization

# Chapter 11

## Gene-centric Annotation

*Bioconductor* provides extensive annotation resources, summarized in Figure 11.1. These can be *gene*-, or *genome*-centric. Annotations can be provided in packages curated by *Bioconductor*, or obtained from web-based resources. Gene-centric *AnnotationDbi* packages include:

- Organism level: e.g. *org.Mm.eg.db*, *Homo.sapiens*.
- Platform level: e.g. *hgu133plus2.db*, *hgu133plus2.probes*, *hgu133plus2.cdf*.
- Homology level: e.g. *hom.Dm.inp.db*.
- System biology level: *GO.db*, *KEGG.db*, *Reactome.db*.

Examples of genome-centric packages include:

- *GenomicFeatures*, to represent genomic features, including constructing reproducible feature or transcript data bases from file or web resources.
- Pre-built transcriptome packages, e.g. *TxDb.Hsapiens.UCSC.hg19.knownGene* based on the *H. sapiens* UCSC hg19 knownGenes track.
- *BSgenome* for whole genome sequence representation and manipulation.
- Pre-built genomes, e.g., *BSgenome.Hsapiens.UCSC.hg19* based on the *H. sapiens* UCSC hg19 build.

Web-based resources include

- *biomaRt* to query *biomart* resource for genes, sequence, SNPs, and etc.
- *rtracklayer* for interfacing with browser tracks, especially the UCSC genome browser.

### 11.1 Gene-centric annotations with *AnnotationDbi*

Organism-level ('org') packages contain mappings between a central identifier (e.g., Entrez gene ids) and other identifiers (e.g. GenBank or Uniprot accession number, RefSeq id, etc.). The name of an org package is always of the form *org.<Sp>.<id>.db* (e.g. *org.Sc.sgd.db*) where *<Sp>* is a 2-letter abbreviation of the organism (e.g. Sc for *Saccharomyces cerevisiae*) and *<id>* is an abbreviation (in lower-case) describing the type of central identifier (e.g. *sgd* for gene identifiers assigned by the *Saccharomyces* Genome Database, or *eg* for Entrez gene ids). The "How to use the '.db' annotation packages" vignette in the *AnnotationDbi* package (org packages are only one type of ".db" annotation packages) is a key reference. The '.db' and most other *Bioconductor* annotation packages are updated every 6 months.

Annotation packages contain an object named after the package itself. These objects are collectively called *AnnotationDb* objects, with more specific classes named *OrgDb*, *ChipDb* or *TranscriptDb* objects. Methods that can be applied to these objects include *cols*, *keys*, *keytypes* and *select*. Common operations for retrieving annotations are summarized in Table 11.1.

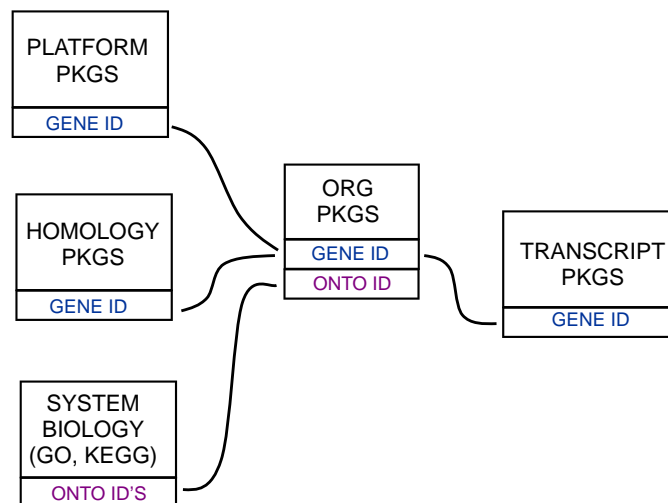


Figure 11.1: Annotation Packages: the big picture

Table 11.1: Common operations for retrieving and manipulating annotations.

Category	Function	Description
Discover	<code>cols</code>	List the kinds of columns that can be returned
	<code>keytypes</code>	List columns that can be used as keys
	<code>keys</code>	List values that can be expected for a given keytype
	<code>select</code>	Retrieve annotations matching <code>keys</code> , <code>keytype</code> and <code>cols</code>
Manipulate	<code>setdiff</code> , <code>union</code> , <code>intersect</code>	Operations on sets
	<code>duplicated</code> , <code>unique</code>	Mark or remove duplicates
	<code>%in%</code> , <code>match</code>	Find matches
	<code>any</code> , <code>all</code>	Are any TRUE? Are all?
	<code>merge</code>	Combine two different <code>data.frames</code> based on shared keys
<i>GRanges</i> *	<code>transcripts</code> , <code>exons</code> , <code>cds</code>	Features (transcripts, exons, coding sequence) as <i>GRanges</i> .
	<code>transcriptsBy</code> , <code>exonsBy</code>	Features group by gene, transcript, etc., as <i>GRangesList</i> .
	<code>cdsBy</code>	



## Exercise 29

What is the name of the *org* package for *Drosophila*? Load it. Display the *OrgDb* object for the [org.Dm.eg.db](#) package. Use the *cols* method to discover which sorts of annotations can be extracted from it.

Use the *keys* method to extract UNIPROT identifiers and then pass those keys in to the *select* method in such a way that you extract the *SYMBOL* (gene symbol) and *KEGG* pathway information for each.

Use *select* to retrieve the *ENTREZ* and *SYMBOL* identifiers of all genes in the *KEGG* pathway 00310.

**Solution:** The *OrgDb* object is named `org.Dm.eg.db`.

```
> library(org.Dm.eg.db)
> cols(org.Dm.eg.db)

[1] "ENTREZID"      "ACCNUM"        "ALIAS"          "CHR"            "CHRLOC"
[6] "CHRLOCEND"     "ENZYME"        "MAP"            "PATH"           "PMID"
[11] "REFSEQ"        "SYMBOL"        "UNIGENE"        "ENSEMBL"        "ENSEMBLPROT"
[16] "ENSEMBLTRANS"  "GENENAME"      "UNIPROT"        "GO"             "EVIDENCE"
[21] "ONTOLOGY"      "GOALL"         "EVIDENCEALL"    "ONTOLOGYALL"    "FLYBASE"
[26] "FLYBASECG"     "FLYBASEPROT"

> keytypes(org.Dm.eg.db)

[1] "ENTREZID"      "ACCNUM"        "ALIAS"          "CHR"            "CHRLOC"
[6] "CHRLOCEND"     "ENZYME"        "MAP"            "PATH"           "PMID"
[11] "REFSEQ"        "SYMBOL"        "UNIGENE"        "ENSEMBL"        "ENSEMBLPROT"
[16] "ENSEMBLTRANS"  "GENENAME"      "UNIPROT"        "GO"             "EVIDENCE"
[21] "ONTOLOGY"      "GOALL"         "EVIDENCEALL"    "ONTOLOGYALL"    "FLYBASE"
[26] "FLYBASECG"     "FLYBASEPROT"

> uniprotKeys <- head(keys(org.Dm.eg.db, keytype="UNIPROT"))
> cols <- c("SYMBOL", "PATH")
> select(org.Dm.eg.db, keys=uniprotKeys, cols=cols, keytype="UNIPROT")

  UNIPROT  SYMBOL  PATH
1  Q8IRZ0  CG3038  <NA>
2  Q95RP8  CG3038  <NA>
3  Q95RU8   G9a    00310
4  Q9W5H1  CG13377  <NA>
5  P39205   cin    <NA>
6  Q24312   ewg    <NA>
```

Selecting UNIPROT and SYMBOL ids of KEGG pathway 00310 is very similar:

```
> kegg <- select(org.Dm.eg.db, "00310", c("UNIPROT", "SYMBOL"), "PATH")
> nrow(kegg)

[1] 36

> head(kegg, 3)

  PATH UNIPROT  SYMBOL
1 00310  Q95RU8   G9a
2 00310  Q9W5E0  Hmt4-20
3 00310  Q9W3N9  CG10932
```

Table 11.2: Selected packages querying web-based annotation services.

Package	Description
<i>biomaRt</i>	<a href="http://biomart.org">http://biomart.org</a> , Ensembl and other annotations
<i>uniprot.ws</i>	<a href="http://uniprot.org">http://uniprot.org</a> , protein annotations
<i>KEGGREST</i>	<a href="http://www.genome.jp/kegg">http://www.genome.jp/kegg</a> , KEGG pathways
<i>SRadb</i>	<a href="http://www.ncbi.nlm.nih.gov/sra">http://www.ncbi.nlm.nih.gov/sra</a> , sequencing experiments.
<i>rtracklayer</i>	<a href="http://genome.ucsc.edu">http://genome.ucsc.edu</a> , genome tracks.
<i>GEOquery</i>	<a href="http://www.ncbi.nlm.nih.gov/geo/">http://www.ncbi.nlm.nih.gov/geo/</a> , array and other data
<i>ArrayExpress</i>	<a href="http://www.ebi.ac.uk/arrayexpress/">http://www.ebi.ac.uk/arrayexpress/</a> , array and other data

### Exercise 30

For convenience, *lrTest*, a *DGEGLM* object from the RNA-seq chapter, is included in the *SequenceAnalysisData* package. The following code loads this data and creates a ‘top table’ of the ten most differentially represented genes. This top table is then coerced to a *data.frame*.

```
> library(SequenceAnalysisData)
> library(edgeR)
> library(org.Dm.eg.db)
> data(lrTest)
> tt <- as.data.frame(topTags(lrTest))
```

Extract the Flybase gene identifiers (*FLYBASE*) from the row names of this table and map them to their corresponding Entrez gene (*ENTREZID*) and symbol ids (*SYMBOL*) using *select*. Use *merge* to add the results of *select* to the top table.

#### Solution:

```
> fbids <- rownames(tt)
> cols <- c("ENTREZID", "SYMBOL")
> anno <- select(org.Dm.eg.db, fbids, cols, "FLYBASE")
> ttanno <- merge(tt, anno, by.x=0, by.y="FLYBASE")
> dim(ttanno)

[1] 10 8

> head(ttanno, 3)

  Row.names logConc logFC LR.statistic PValue   FDR ENTREZID SYMBOL
1 FBgn0000071    -11  2.8         183 1.1e-41 1.1e-38  40831   Ama
2 FBgn0024288    -12 -4.7         179 7.1e-41 6.3e-38  45039 Sox100B
3 FBgn0033764    -12  3.5         188 6.8e-43 7.8e-40   <NA>   <NA>
```

## 11.2 *biomaRt* and other web-based resources

A short summary of select *Bioconductor* packages enabling web-based queries is in Table 11.2.

### 11.2.1 Using *biomaRt*

The *biomaRt* package offers access to the online *biomart* resource. this consists of several data base resources, referred to as ‘marts’. Each mart allows access to multiple data sets; the *biomaRt* package provides methods for mart and data set discovery, and a standard method *getBM* to retrieve data.

### Exercise 31

Load the *biomaRt* package and list the available marts. Choose the ensembl mart and list the datasets for that mart. Set up a mart to use the ensembl mart and the *hsapiens\_gene\_ensembl* dataset.

A *biomaRt* dataset can be accessed via *getBM*. In addition to the mart to be accessed, this function takes filters and attributes as arguments. Use *filterOptions* and *listAttributes* to discover values for these arguments. Call *getBM* using filters and attributes of your choosing.

### Solution:

```
> library(biomaRt)
> head(listMarts(), 3)                ## list the marts
> head(listDatasets(useMart("ensembl")), 3) ## mart datasets
> ensembl <-                          ## fully specified mart
+   useMart("ensembl", dataset = "hsapiens_gene_ensembl")
> head(listFilters(ensembl), 3)       ## filters
> myFilter <- "chromosome_name"
> head(filterOptions(myFilter, ensembl), 3) ## return values
> myValues <- c("21", "22")
> head(listAttributes(ensembl), 3)    ## attributes
> myAttributes <- c("ensembl_gene_id", "chromosome_name")
> ## assemble and query the mart
> res <- getBM(attributes = myAttributes, filters = myFilter,
+             values = myValues, mart = ensembl)
```

Use *head(res)* to see the results.

# Chapter 12

## Genomic Annotation

### 12.1 Whole genome sequences

There are a diversity of packages and classes available for representing large genomes. Several include:

***TxDb.\**** For transcript and other genome / coordinate annotation.

***BSgenome*** For whole-genome representation. See `available.packages` for pre-packaged genomes, and the vignette ‘How to forge a BSgenome data package’ in the

***Homo.sapiens*** For integrating *TxDb.\** and *org.\** packages.

***SNPlocs.\**** For model organism SNP locations derived from dbSNP.

**FaFile** (*Rsamtools*) for accessing indexed FASTA files.

***SIFT.\**, *PolyPhen*** Variant effect scores.

### 12.2 Gene models

#### 12.2.1 *TxDb.\** packages for model organisms

Genome-centric packages are very useful for annotations involving genomic coordinates. It is straight-forward, for instance, to discover the coordinates of coding sequences in regions of interest, and from these retrieve corresponding DNA or protein coding sequences. Other examples of the types of operations that are easy to perform with genome-centric annotations include defining regions of interest for counting aligned reads in RNA-seq experiments and retrieving DNA sequences underlying regions of interest in ChIP-seq analysis, e.g., for motif characterization.




Figure 12.1: Gene model showing exons on three transcripts.

### Exercise 32

Load the ‘transcript.db’ package relevant to the dm3 build of *D. melanogaster*. Use `select` and `friends` to select the Flybase gene ids of the top table `tt` and the Flybase transcript names (`TXNAME`) and Entrez gene identifiers (`GENEID`).

Use `cdsBy` to extract all coding sequences, grouped by transcript. Subset the coding sequences to contain just the transcripts relevant to the top table. How many transcripts are there? What is the structure of the first transcript’s coding sequence?

Load the ‘BSgenome’ package for the dm3 build of *D. melanogaster*. Use the coding sequences ranges of the previous part of this exercise to extract the underlying DNA sequence, using the `extractTranscriptsFromGenome` function. Use `Biostrings`’ `translate` to convert DNA to amino acid sequences.

**Solution:** The following loads the relevant Transcript.db package, and creates a more convenient alias to the *TranscriptDb* instance defined in the package.

```
> library(TxDb.Dmelanogaster.UCSC.dm3.ensGene)
> txdb <- TxDb.Dmelanogaster.UCSC.dm3.ensGene
```

We also need the data – flybase IDs from our differential expression analysis.

```
> library(SequenceAnalysisData)
> data(lrTest)
> fbids <- rownames(topTags(lrTest))
```

We can discover available keys (using `keys`) and columns (`cols`) in `txdb`, and then use `select` to retrieve the transcripts associated with each differentially expressed gene. The mapping between gene and transcript is not one-to-one – some genes have more than one transcript.

```
> txnm <- select(txdb, fbids, "TXNAME", "GENEID")
> nrow(txnm)
```

```
[1] 19
```

```
> head(txnm, 3)
```

	GENEID	TXNAME
1	FBgn0039155	FBtr0084549
2	FBgn0039827	FBtr0085755
3	FBgn0039827	FBtr0085756

The *TranscriptDb* instances can be queried for data that is more structured than simple data frames, and in particular return *GRanges* or *GRangesList* instances to represent genomic coordinates. These queries are performed using `cdsBy` (coding sequence), `transcriptsBy` (transcripts), etc., where a function argument `by` specifies how coding sequences or transcripts are grouped. Here we extract the coding sequences grouped by transcript, returning the transcript names, and subset the resulting *GRangesList* to contain just the transcripts of interest to us. The first transcript is composed of 6 distinct coding sequence regions.

```
> cds <- cdsBy(txdb, "tx", use.names=TRUE)[txnm$TXNAME]
> length(cds)
```

```
[1] 19
```

```
> cds[1]
```

```
GRangesList of length 1:
```

```
$FBtr0084549
```

```
GRanges with 6 ranges and 3 metadata columns:
```

	seqnames	ranges	strand	cds_id	cds_name	exon_rank
	<Rle>	<IRanges>	<Rle>	<integer>	<character>	<integer>
[1]	chr3R	[19970946, 19971592]	+	39378	<NA>	2
[2]	chr3R	[19971652, 19971770]	+	39379	<NA>	3
[3]	chr3R	[19971831, 19972024]	+	39380	<NA>	4
[4]	chr3R	[19972088, 19972461]	+	39381	<NA>	5
[5]	chr3R	[19972523, 19972589]	+	39382	<NA>	6
[6]	chr3R	[19972918, 19973094]	+	39383	<NA>	7

---

seqlengths:

chr2L	chr2R	chr3L	chr3R	...	chrXHet	chrYHet	chrUextra
23011544	21146708	24543557	27905053	...	204112	347038	29004656

The following code loads the appropriate BSgenome package; the *Dmelanogaster* object refers to the whole genome sequence represented in this package. The remaining steps extract the DNA sequence of each transcript, and translates these to amino acid sequences. Issues of strand are handled correctly.

```
> library(BSgenome.Dmelanogaster.UCSC.dm3)
> txx <- extractTranscriptsFromGenome(Dmelanogaster, cds)
> length(txx)

[1] 19

> head(txx, 3)

A DNAStringSet instance of length 3
width seq names
[1] 1578 ATGGGCAGCATGCAAGTGGCGCT...TGCAGATCAAGTGCAGCGACTAG FBtr0084549
[2] 2760 ATGCTGCGTTATCTGGCGCTTTC...TTGCTGCCCCATTGGAACCTTTAG FBtr0085755
[3] 2217 ATGGCACTCAAGTTTCCACAGT...TTGCTGCCCCATTGGAACCTTTAG FBtr0085756

> head(translate(txx), 3)

A AAStringSet instance of length 3
width seq
[1] 526 MGSMQVALLALLVLGQLFPSAVANGSSSYSTST...VLDDSRNVFTFTTPKCENFRKRFPKLQIKCSD*
[2] 920 MLRYLALSEAGIAKLPRPQSRCHSEKGVWGYKP...YCGRCEAPTPATGIGKVHKREVDEIVAAPFEL*
[3] 739 MALKFPTVKRYGGEAESMLAFFWQLRDSVQAN...YCGRCEAPTPATGIGKVHKREVDEIVAAPFEL*
```

## 12.3 UCSC tracks

### 12.3.1 Easily creating *TranscriptDb* objects from GTF files

#### Exercise 33

(Adapted from<sup>1</sup>) The ‘track’ curated annotations at UCSC are a great resource; this exercise creates a *TranscriptDb* instance from one such track.

- Load the *GenomicFeatures* and *rtracklayer* packages.
- Discover available genomes with *ucscGenomes*, and available tables with the *supportedUCSCTables*.

<sup>1</sup><http://www.sph.emory.edu/~hwu/teaching/bioc/bios560R.html>

- c. Use the `makeTranscriptDbFromUCSC` from a suitable track, e.g., genome `genome="ce2"`, `tablename="refGene"`. There are some warnings; are these something to be concerned about?
- d. Exercise the object that you've created, e.g., exploring the basis of the warnings.
- e. Save and load the `TranscriptDb` object you've created, illustrating how one can make these annotations convenient and reproducible.
- f. What's the difference between `makeTranscriptDbFromUCSC` and `makeFeatureDbFromUCSC`? Where else can transcript and feature data bases be made from?

**Solution:** Load the `GenomicFeatures` package and discover available genomes and tables:

```
> library(rtracklayer)
> library(GenomicFeatures)
> ## genomes
> gnms <- ucscGenomes()
> nrow(gnms)

[1] 145

> gnms[grep("elegans", gnms$species),]

      db species      date      name
134 ce10 C. elegans Oct. 2010 WormBase v. WS220
135 ce6 C. elegans May 2008 WormBase v. WS190
136 ce4 C. elegans Jan. 2007 WormBase v. WS170
137 ce2 C. elegans Mar. 2004 WormBase v. WS120

> ## tables
> tbls <- supportedUCSCTables()
> nrow(tbls)

[1] 25

> head(tbls)

      track      subtrack
knownGene      UCSC Genes      <NA>
knownGeneOld3  Old UCSC Genes      <NA>
wgEncodeGencodeManualV3  Gencode Genes  Gencode Manual
wgEncodeGencodeAutoV3    Gencode Genes  Gencode Auto
wgEncodeGencodePolyaV3   Gencode Genes  Gencode PolyA
ccdsGene           CCDS           <NA>
```

Make the `TranscriptDb` object (this will take a minute)

```
> ## Not run
> txdb <- makeTranscriptDbFromUCSC("ce10", "refGene")
> saveDb(txdb, file="/path/to/file.sqlite")
```

The warnings during object creation are about unusual lengths for CDS (coding sequences should be in multiples of 3, since there are three nucleotide residues per amino acid residue).

```
1: In .extractUCSCCdsStartEnd(cdsStart[i], cdsEnd[i], ... :
UCSC data anomaly in transcript NM_001129046: the cds cumulative
length is not a multiple of 3
```

but we seem to have a useful object with relevant metadata information for reproducible research:

```
> txdb

TranscriptDb object:
| Db type: TranscriptDb
| Supporting package: GenomicFeatures
| Data source: UCSC
| Genome: ce10
| Organism: Caenorhabditis elegans
| UCSC Table: refGene
| Resource URL: http://genome.ucsc.edu/
| Type of Gene ID: Entrez Gene ID
| Full dataset: yes
| miRBase build ID: NA
| transcript_nrow: 48714
| exon_nrow: 152542
| cds_nrow: 129947
| Db created by: GenomicFeatures package from Bioconductor
| Creation time: 2013-02-10 10:33:50 -0800 (Sun, 10 Feb 2013)
| GenomicFeatures version at creation time: 1.11.8
| RSQLite version at creation time: 0.11.2
| DBSCHEMAVERSION: 1.0
```

Let's investigate the source of the warnings – what fraction of the CDS have lengths that are not a multiple of 3? To do this we'll need to assemble the figure out the sum of the widths of the coding sequence exons in each transcript. Start by extracting all coding sequence exons, grouped by transcript; verify that this is a *GRangesList* with a reasonable number of entries.

```
> cdsByTx <- cdsBy(txdb, "tx", use.names=TRUE)
> length(cdsByTx)
```

```
[1] 26146
```

```
> cdsByTx[1:2]
```

GRangesList of length 2:

\$NM\_058259

GRanges with 3 ranges and 3 metadata columns:

	seqnames	ranges	strand	cds_id	cds_name	exon_rank
	<Rle>	<IRanges>	<Rle>	<integer>	<character>	<integer>
[1]	chrI	[11641, 11689]	+	1	<NA>	1
[2]	chrI	[14951, 15160]	+	2	<NA>	2
[3]	chrI	[16473, 16585]	+	3	<NA>	3

\$NM\_058264

GRanges with 5 ranges and 3 metadata columns:

	seqnames	ranges	strand	cds_id	cds_name	exon_rank
[1]	chrI	[43733, 43961]	+	4	<NA>	1
[2]	chrI	[44030, 44234]	+	5	<NA>	2
[3]	chrI	[44281, 44324]	+	6	<NA>	3
[4]	chrI	[44372, 44468]	+	7	<NA>	4
[5]	chrI	[44521, 44677]	+	8	<NA>	5



```

---
seqlengths:
  chrI    chrII    chrIII    chrIV    chrV    chrX    chrM
15072423 15279345 13783700 17493793 20924149 17718866 13794

```

Extract the width of each exon; this is an *IntegerList* instance with one element of the list for each transcript, and one integer value for each exon.

```

> wd <- width(cdsByTx)
> length(wd)

[1] 26146

> head(wd, 3)

CompressedIntegerList of length 3
[["NM_058259"]] 49 210 113
[["NM_058264"]] 229 205 44 97 157
[["NM_001026606"]] 139 163 183 166

```

Use `sum` to add up the widths within each list element; note that we're using the `sum,CompressedIntegerList-method`, and that this has been specialized to do the summation within list elements.

```

> head(sum(wd))

      NM_058259    NM_058264 NM_001026606    NM_058265 NM_001026607    NM_182094
      372          732          651          1341          1620          1221

```

We now have a standard *R* vector; a one-liner asking about the number of transcripts with exons that do not sum to 3 is

```

> table( (sum(width(cdsBy(txdb, "tx")))) %% 3) != 0 )

FALSE  TRUE
25676   470

```

### Exercise 34

(Adapted from<sup>2</sup>) Suppose you have a list of transcription factor binding sites on *hg19*. How would you obtain (a) the GC content of each site and (b) the percentage of gene promoters covered by the binding sites?

**Solution:** As an outline of a solution, the steps for calculating GC content might be

- Represent the list of transcription factor binding sites ('regions of interest') as a *GRanges* instance, `roi`.
- Load *BSgenome* and the appropriate genome package, e.g., *BSgenome.Hsapiens.UCSC.hg19*.
- Use `getSeq` to retrieve the sequences, `seqs <- getSeq(Hsapiens, roi)`.
- Use `alphabetFrequency(seqs)` to summarize nucleotide use, and simple *R* functions to determine GC content of each region of interest.
- Summarize these as density plots, etc. A meaningful extension of this exercise might compare the observed GC content to the expected content, where expectation is the product of the independent G and C frequencies.

---

<sup>2</sup><http://www.sph.emory.edu/~hwu/teaching/bioc/bios560R.html>

To calculate the percentage of promoters covered by binding sites, we might

- a. Load the reference genome *TxDb* package, [\*TxDb.Hsapiens.UCSC.hg19.knownGene\*](#).
- b. Query the package for promoters using the `promoters` function, or otherwise manipulating exon or transcript coordinates to get a *GRanges* or *GRangesList* representing genomic regions of interest, `groi`.
- c. Use `countOverlaps(groi, roi)` to find how many transcription factor binding sites overlap each promoter, and from there use standard *R* functions to tally the number of promoters that have zero overlaps.

## Chapter 13

# Visualizing Sequence Data

*R* has some great visualization packages; essential references include [5] for a general introduction, Murrell [17] for base graphics, Sarkar [23] for *lattice*, and Wickham [25] for *ggplot2*. Here we take a quick tour of visualization facilities tailed for sequence data and using *Bioconductor* approaches.

### 13.1 *Gviz*

The *Gviz* package produces very elegant data organized in a more-or-less familiar ‘track’ format. The following exercises walk through the *Gviz User guide* Section 2.

#### Exercise 35

Load the *Gviz* package and sample *GRanges* containing genomic coordinates of CpG islands. Create a couple of variables with information on the chromosome and genome of the data (how can this information be extracted from the *cpgIslands* object?).

```
> library(Gviz)
> data(cpgIslands)
> chr <- "chr7"
> genome <- "hg19"
```

The basic idea is to create a track, perhaps with additional attributes, and to plot it. There are different types of track, and we create these one at a time. We start with a simple annotation track

```
> atrack <- AnnotationTrack(cpgIslands, name="CpG")
> plotTracks(atriack)
```

Then add a track that represents genomic coordinates. Tracks are combined during when plotted, as a simple list. The vertical ordering of tracks is determined by their position in the list.

```
> gtrack <- GenomeAxisTrack()
> plotTracks(list(gtrack, atrack))
```

We can add an ideogram to provide overall orientation...

```
> itrack <- IdeogramTrack(genome=genome, chromosome=chr)
> plotTracks(list(itrack, gtrack, atrack))
```

and a more elaborate gene model, as an *data.frame* or *GRanges* object with specific columns of metadata.

```

> data(geneModels)
> grtrack <-
+   GeneRegionTrack(geneModels, genome=genome,
+                   chromosome=chr, name="Gene Model")
> tracks <- list(itrack, gtrack, atrack, grtrack)
> plotTracks(tracks)

```

Zooming out changes the location box on the ideogram

```

> plotTracks(tracks, from=2.5e7, to=2.8e7)

```

When zoomed in we can add sequence data

```

> library(BSgenome.Hsapiens.UCSC.hg19)
> strack <- SequenceTrack(Hsapiens, chromosome=chr)
> plotTracks(c(tracks, strack), from=26450430, to=26450490, cex=.8)

```

As the *Gviz* vignette humbly says, ‘so far we have replicated the features of a whole bunch of other genome browser tools out there’. We’d like to be able integrate our data into these plots, with a rich range of plotting options. The key is the `DataTrack` function, which we demonstrate with some simulated data; this final result is shown in Figure 13.1.

```

> ## some data
> lim <- c(26700000, 26900000)
> coords <- seq(lim[1], lim[2], 101)
> dat <- runif(length(coords) - 1, min=-10, max=10)
> ## DataTrack
> dtrack <-
+   DataTrack(data=dat, start=coords[-length(coords)],
+             end= coords[-1], chromosome=chr, genome=genome,
+             name="Uniform Random")
> plotTracks(c(tracks, dtrack))

```

Section 4.3 of the *Gviz* vignette illustrates flexibility of the data track.

## 13.2 *ggbio*

The *ggbio* package complements facilities in *Gviz*. It uses the central metaphors of the ‘grammar of graphics’ made popular by the *ggplot2* package, and integrates closely with the *Bioconductor* ranges infrastructure. In addition to the grammar of graphics approach, the package offers a wider range of plots, for instance circular plots. The use of the package is covered in its vignette.


## 13.3 *shiny* for easy interactive reports

As a final example of visualization, the *shiny* package and web site <sup>1</sup> has recently been introduced. It offers a new model for developing interactive, browser-based visualizations. These visualizations could be an excellent way to provide sophisticated exploratory or summary analysis in a very accessible way. The idea is to write a ‘user interface’ component that describes how a page is to be presented to users, and a ‘server’ that describes how the data are to be calculated or modified in responses to user choices. The programming model is ‘reactive’, where changes in a user choice automatically trigger re-calculations in the server. This reactive model is like in a spreadsheet with a formula, where adjusting a cell that the formula references

---

<sup>1</sup><http://www.rstudio.com/shiny/>

Figure 13.1: *Gviz* ideogram, genome coordinate, annotation, and data tracks.



triggers re-calculation of the formula. Just like in a spreadsheet, someone creating a *shiny* application does not have to work hard to make reactivity work.

**WARNING:** The following demos were only available during the course.

There are two short demos available for use during the course. The first

```
> source("http://bioconductor.org/scratch-repos/pkgInstall.R")
> demo1()
```

Uses a *SummarizedExperiment* object to contain the data and results of the *DESeq* work flow from yesterday, coupled with the annotation resources that we explored today. The user can choose genes to display based on p-value and log-fold change in the ‘top table’ of genes. In reaction to these choices, the heat map on one tab and annotations on a second, are updated.

One of the nice features of *shiny* is it separates responsibility for doing manipulations of the data (*R*’s responsibility) from obligations to display it (*javascript*’s responsibility). This means that it can be quite ‘easy’ to incorporate complicate visualizations, as illustrated in

```
> demo2()
```

This represents a gene network under various conditions. The data is in *R*, and changing the slider causes *R* to update the data, but the visualization uses the javascript development version of *cytoscape*. This offers great opportunities for interaction with different projects.

# References

- [1] S. Anders and W. Huber. Differential expression analysis for sequence count data. *Genome Biol*, 11(10):R106, 2010.
- [2] D. Bentley, S. Balasubramanian, H. Swerdlow, G. Smith, J. Milton, C. Brown, K. Hall, D. Evers, C. Barnes, H. Bignell, et al. Accurate whole human genome sequencing using reversible terminator chemistry. *Nature*, 456(7218):53–59, 2008.
- [3] A. N. Brooks, L. Yang, M. O. Duff, K. D. Hansen, J. W. Park, S. Dudoit, S. E. Brenner, and B. R. Graveley. Conservation of an RNA regulatory map between *Drosophila* and mammals. *Genome Research*, pages 193–202, 2011.
- [4] J. H. Bullard, E. Purdom, K. D. Hansen, and S. Dudoit. Evaluation of statistical methods for normalization and differential expression in mrna-seq experiments. *BMC bioinformatics*, 11(1):94, 2010.
- [5] W. Chang. *R Graphics Cookbook*. O’Reilly Media, Incorporated, 2012.
- [6] P. Dalgaard. *Introductory Statistics with R*. Springer, 2nd edition, 2008.
- [7] R. Gentleman. *R Programming for Bioinformatics*. Computer Science & Data Analysis. Chapman & Hall/CRC, Boca Raton, FL, 2008.
- [8] F. Hahne, W. Huber, R. Gentleman, and S. Falcon. *Bioconductor case studies*. Springer, 2008.
- [9] K. D. Hansen, S. E. Brenner, and S. Dudoit. Biases in illumina transcriptome sequencing caused by random hexamer priming. *Nucleic acids research*, 38(12):e131–e131, 2010.
- [10] K. D. Hansen, R. A. Irizarry, and Z. Wu. Removing technical variability in RNA-seq data using conditional quantile normalization. *Biostatistics*, 13(2):204–216, 2012.
- [11] B. Langmead, C. Trapnell, M. Pop, and S. L. Salzberg. Ultrafast and memory-efficient alignment of short DNA sequences to the human genome. *Genome Biol.*, 10:R25, 2009.
- [12] H. Li and R. Durbin. Fast and accurate long-read alignment with Burrows-Wheeler transform. *Bioinformatics*, 26:589–595, Mar 2010.
- [13] J. C. Marioni, C. E. Mason, S. M. Mane, M. Stephens, and Y. Gilad. Rna-seq: an assessment of technical reproducibility and comparison with gene expression arrays. *Genome research*, 18(9):1509–1517, 2008.
- [14] N. Matloff. *The Art of R Programming*. No Starch Press, 2011.
- [15] J. Meys and A. de Vries. *R For Dummies*. For Dummies, 2012.
- [16] A. Mortazavi, B. A. Williams, K. McCue, L. Schaeffer, and B. Wold. Mapping and quantifying mammalian transcriptomes by rna-seq. *Nature methods*, 5(7):621–628, 2008.
- [17] P. Murrell. *R graphics*. Chapman & Hall/CRC, 2005.

- [18] P. J. Park. ChIP-seq: advantages and challenges of a maturing technology. *Nat. Rev. Genet.*, 10:669–680, Oct 2009. [PubMed Central:[PMC3191340](#)] [DOI:[10.1038/nrg2641](#)] [PubMed:[19736561](#)].
- [19] R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2013. ISBN 3-900051-07-0.
- [20] M. D. Robinson, D. J. McCarthy, and G. K. Smyth. edgeR: a Bioconductor package for differential expression analysis of digital gene expression data. *Bioinformatics*, 26:139–140, Jan 2010.
- [21] C. Ross-Innes, R. Stark, A. Teschendorff, K. Holmes, H. Ali, M. Dunning, G. Brown, O. Gojis, I. Ellis, A. Green, et al. Differential oestrogen receptor binding is associated with clinical outcome in breast cancer. *Nature*, 481(7381):389–393, 2012.
- [22] J. Rothberg and J. Leamon. The development and impact of 454 sequencing. *Nature biotechnology*, 26(10):1117–1124, 2008.
- [23] D. Sarkar. *Lattice: multivariate data visualization with R*. Springer, 2008.
- [24] M. Schmidberger, M. Morgan, D. Eddelbuettel, H. Yu, L. Tierney, and U. Mansmann. State of the art in parallel computing with r. *Journal of Statistical Software*, 31(1):1–27, 8 2009.
- [25] H. Wickham. *ggplot2: elegant graphics for data analysis*. Springer Publishing Company, Incorporated, 2009.
- [26] H. Wu, C. Wang, and Z. Wu. A new shrinkage estimator for dispersion improves differential expression detection in rna-seq data. *Biostatistics*, 2012.

# Part V

## Appendix



## Appendix A

### *DESeq* vignette

## Appendix B

### *VariantTools* vignette